

Delta and Databricks vs SQL Server

Who am I



- 18+ Years Experience
- Data Engineering & AI Consultant
- Intensive Software & Data Engineering Experience
- Microsoft AI MVP
- Public Speaker
- Community Organiser



X @annawykes

in @annawykes



Who am I



- SQL Server DBA/Developer
- 20 years + experience
- Microsoft Data Platform MVP
- User Group Leader
- SQLBits Organiser



Who am I



- 9+ Years Experience
- Database Engineer
- Public Speaker
- Data Bristol Organiser
- Pydata Bristol Organiser



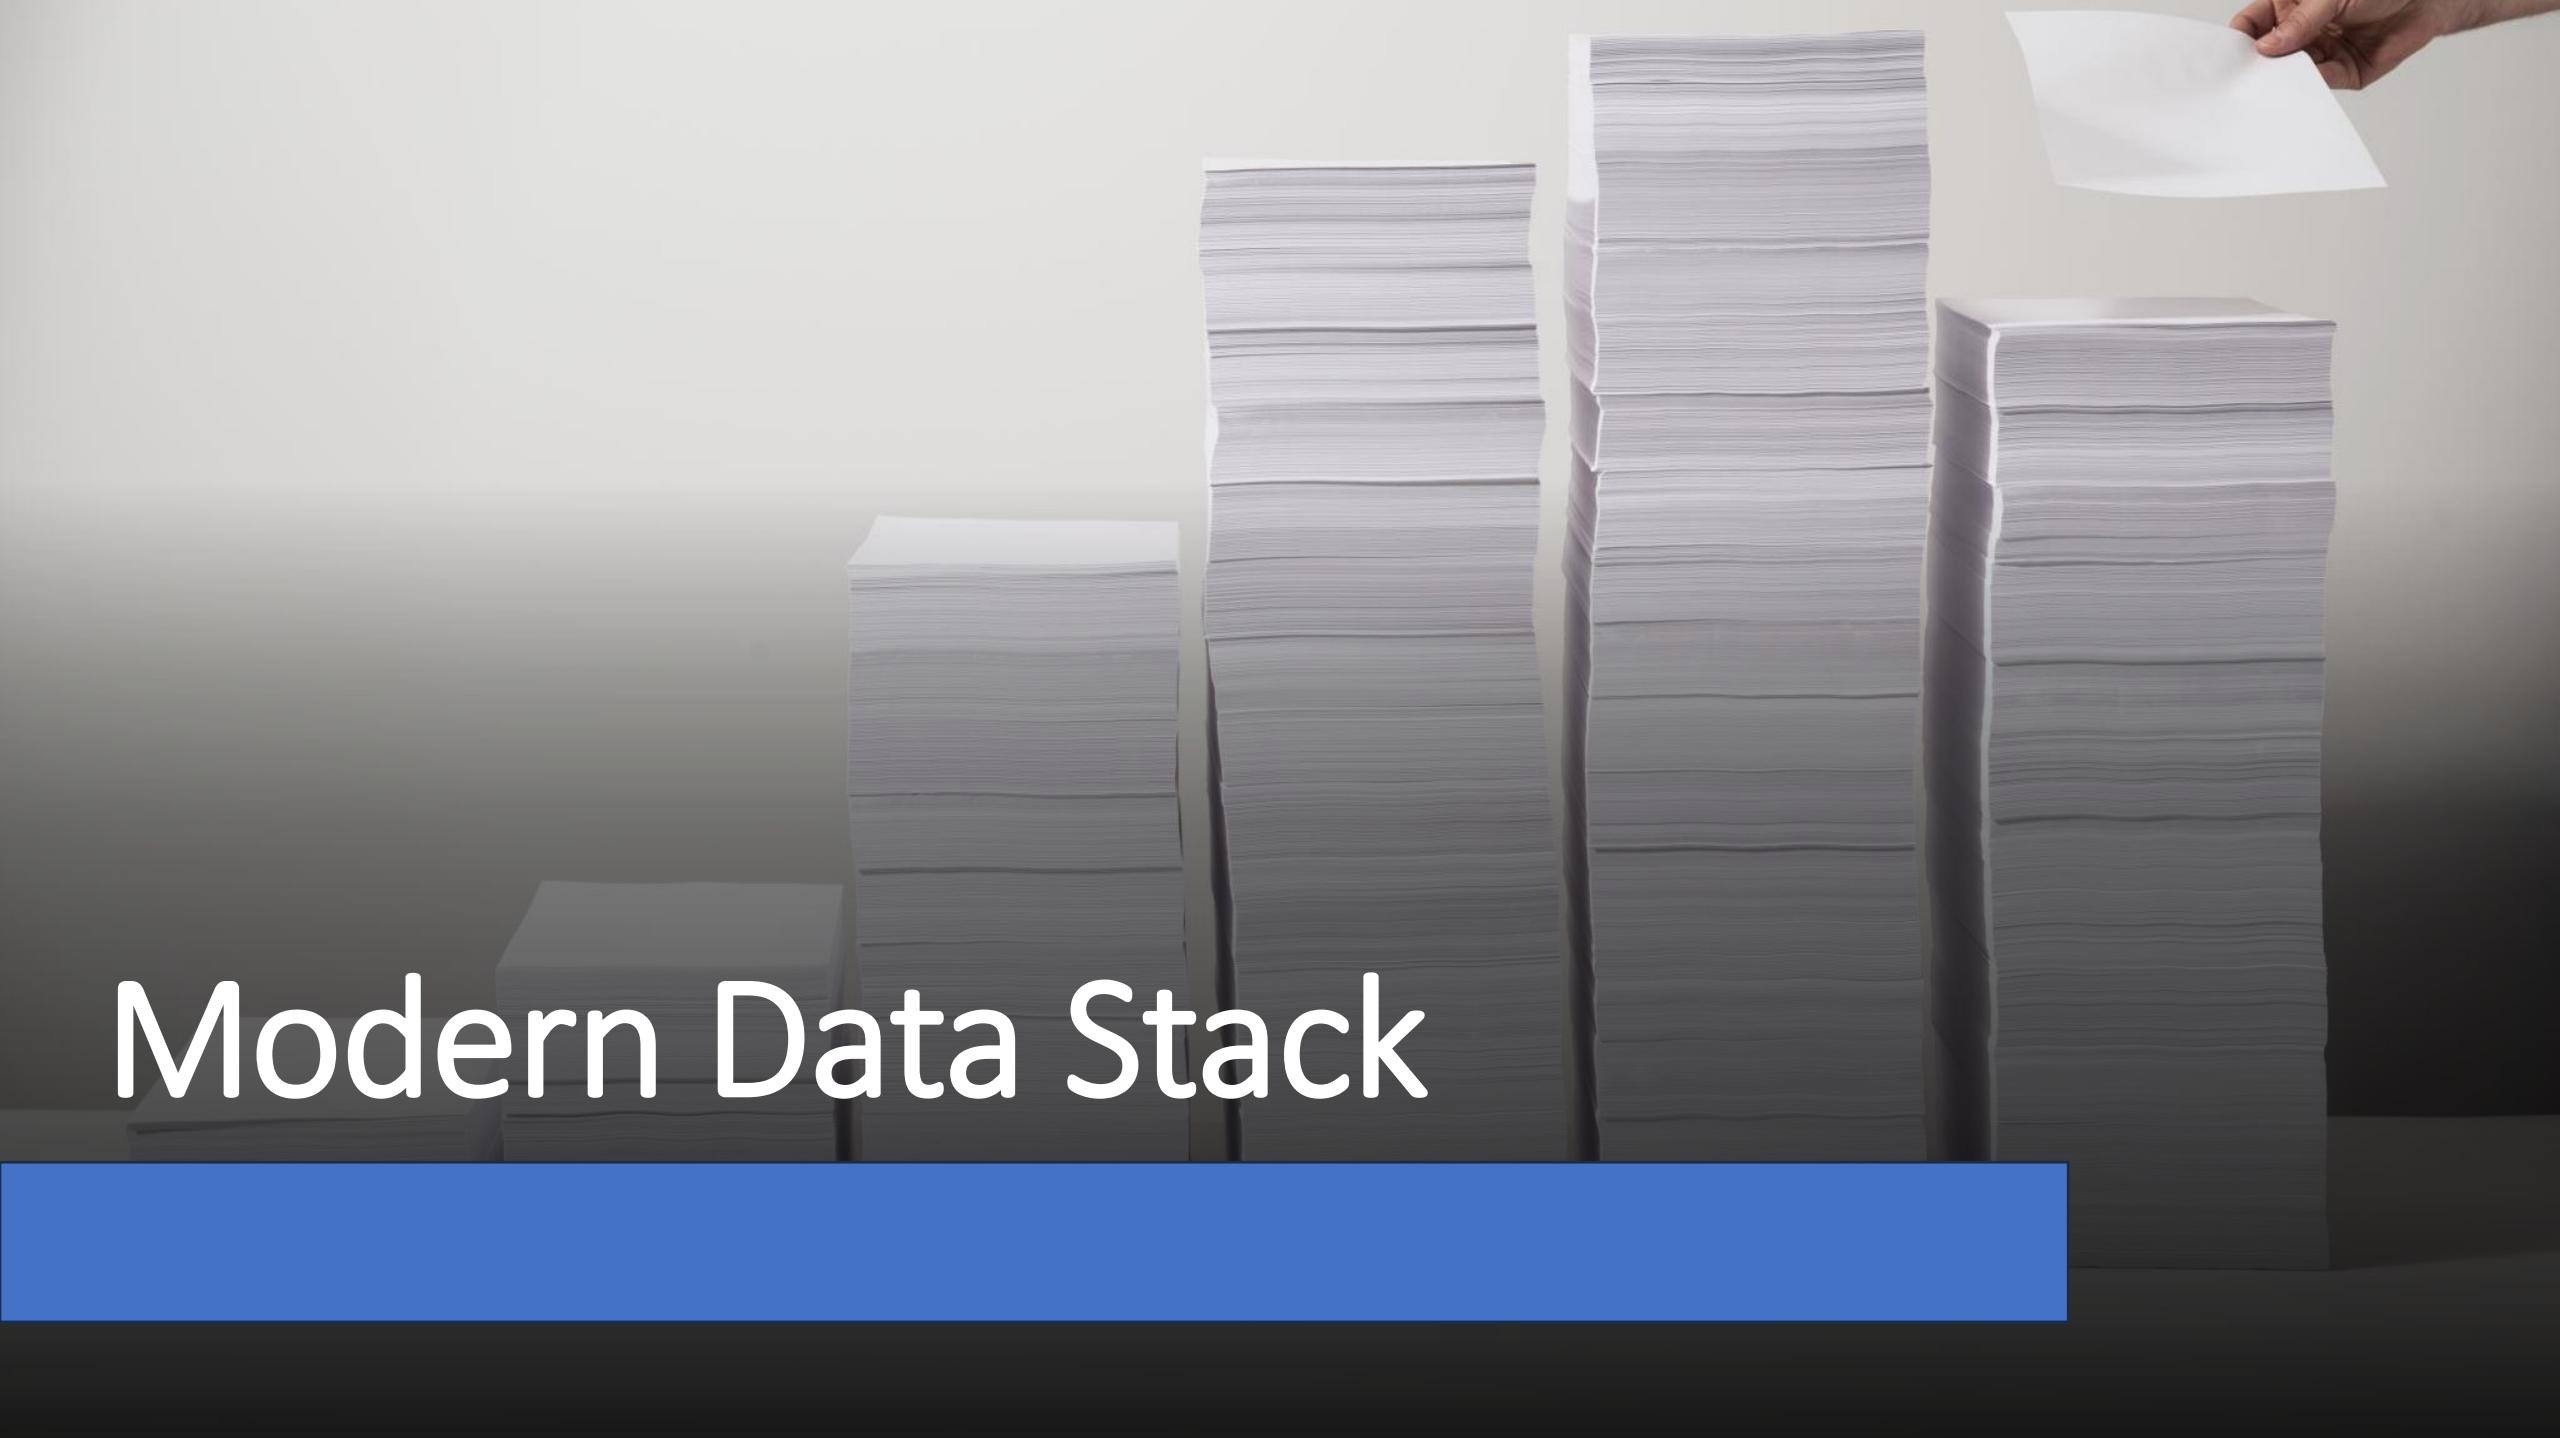
 james-c-yarrow

Why?

Agenda

- Modern Data Stack and where we fit
- Databricks Clusters vs SQL Engine
- ACID Transactions
- Parquet/Delta Storage vs Storage Structures
- Partitioning Delta vs SQL Server
- Z-ordering and Liquid Clustering vs Index Maintenance
- Change Data Capture (CDC)
- Delta History and Time Travel vs SQL Server Temporal Tables
- Delta Vacuum vs Data Retention
- Identity Columns, Primary & Foreign Keys
- Security (Unity Catalog) vs SQL Server Security

Modern Data Stack

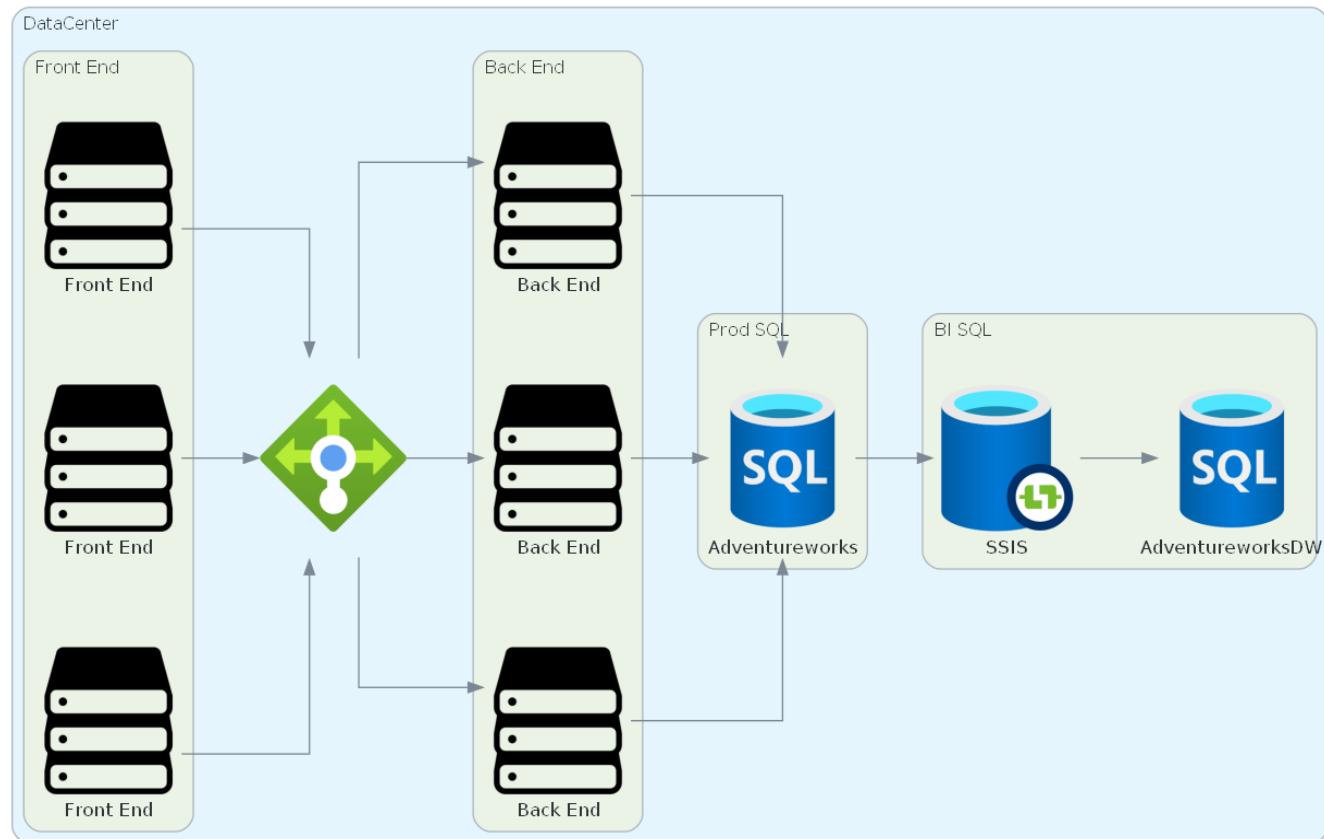


Modern Data Stack

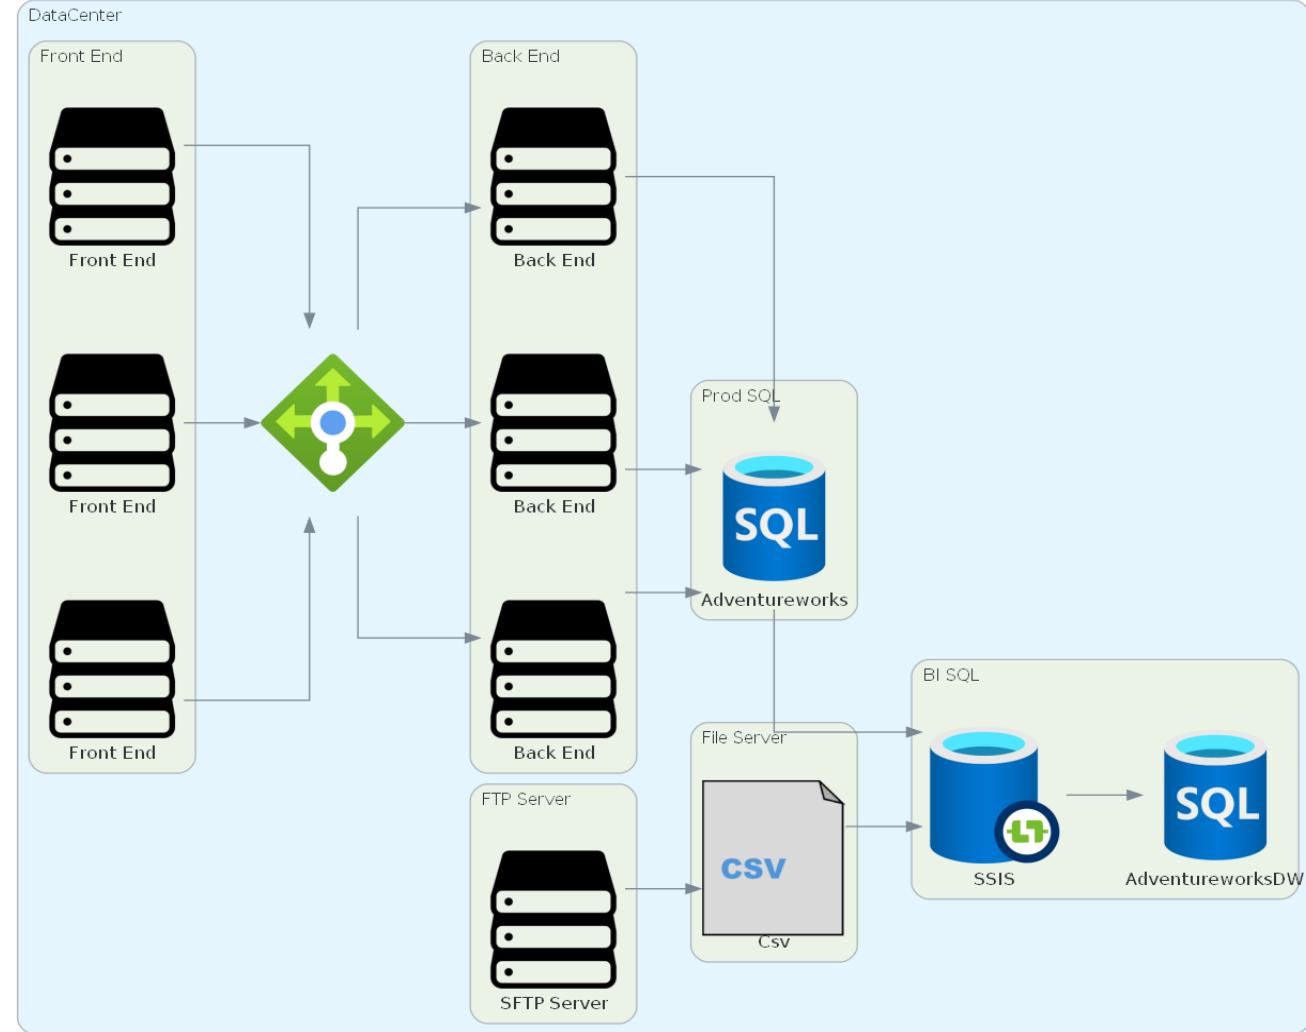
- The same data stack with modern components and the capabilities they bring
- Same data model challenges
- More Distribution
- More Centralisation
- Let take a quick journey from our traditional SQL Server to SQL Server ETL



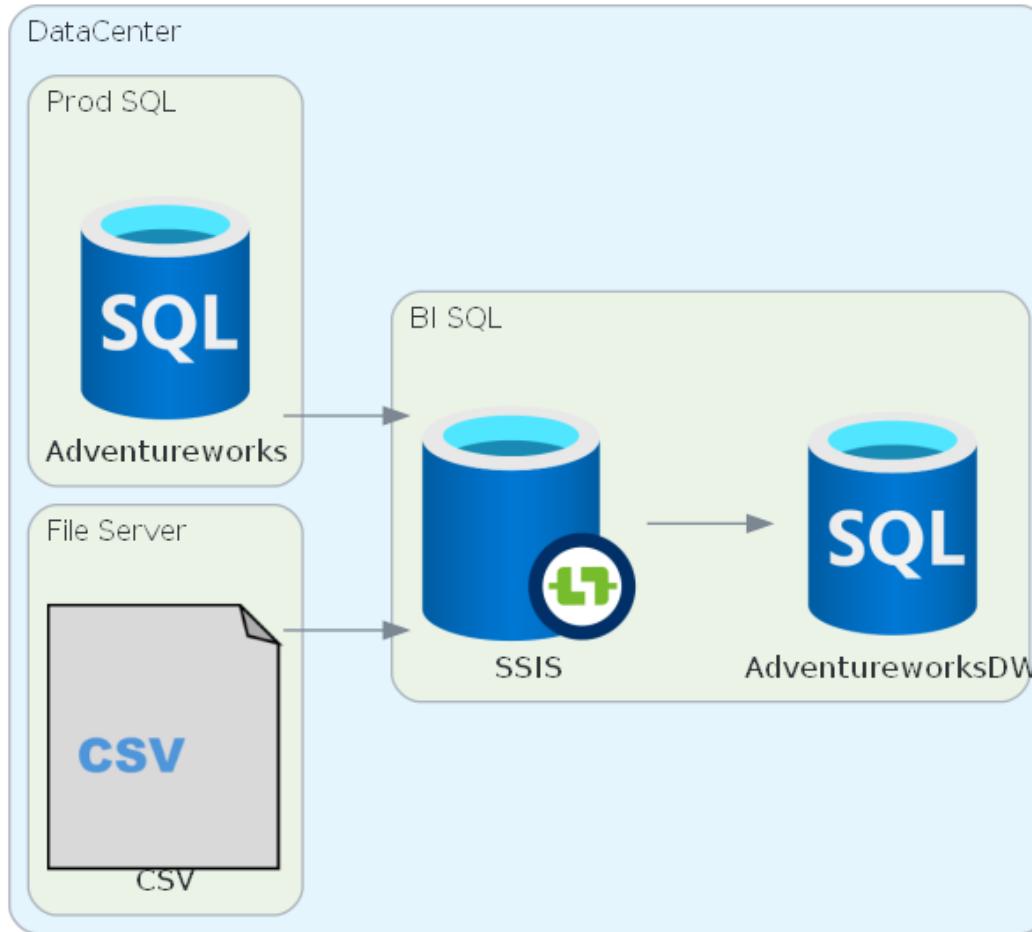
Modern Data Stack



Modern Data Stack

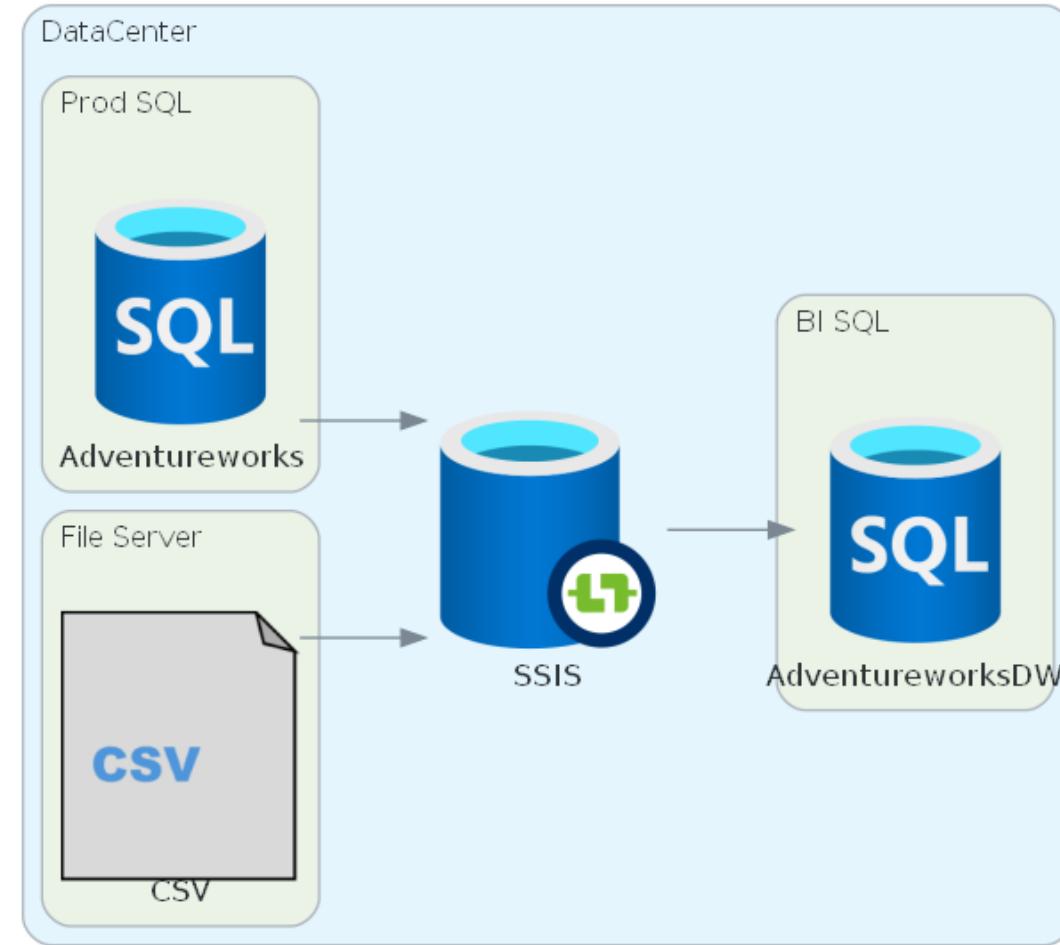


Modern Data Stack

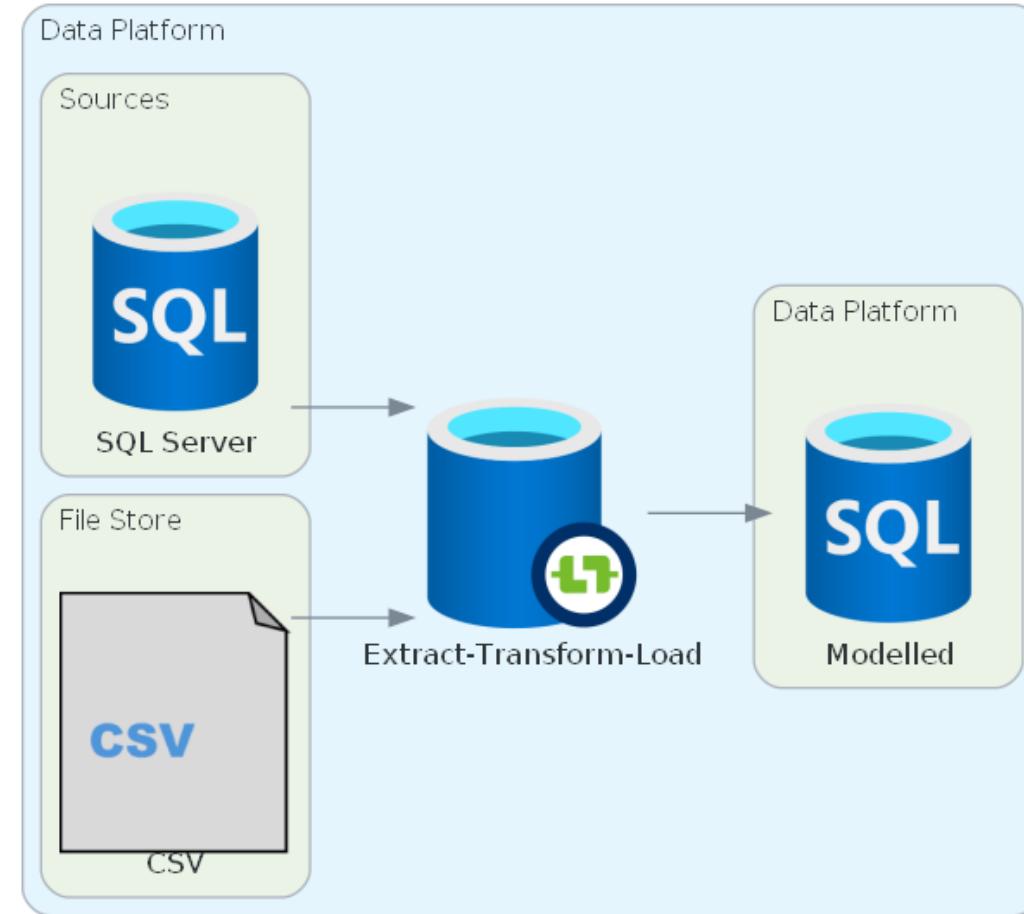


Modern Data Stack

SSIS does a lot!

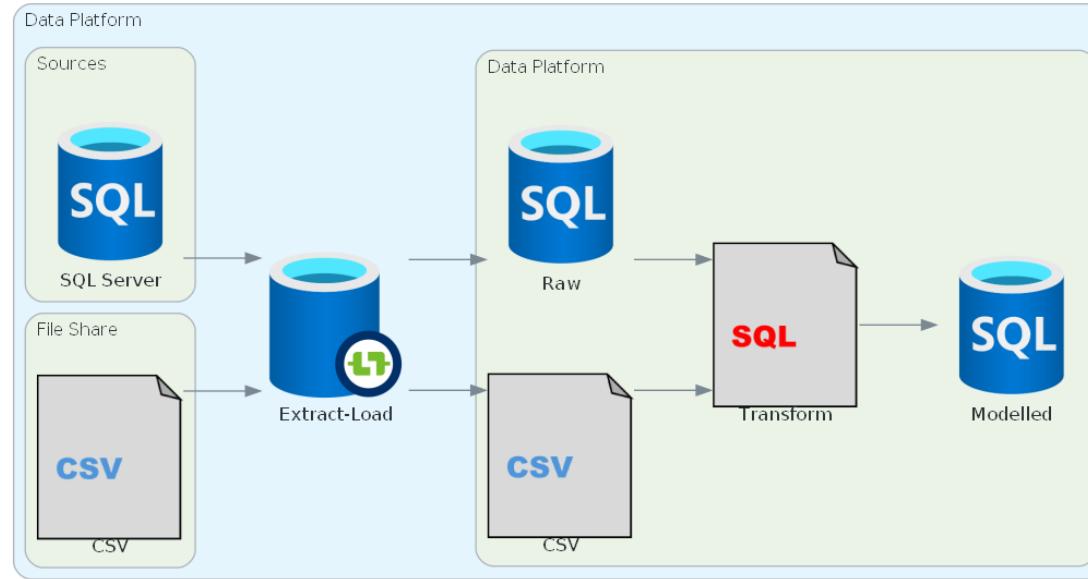


Modern Data Stack



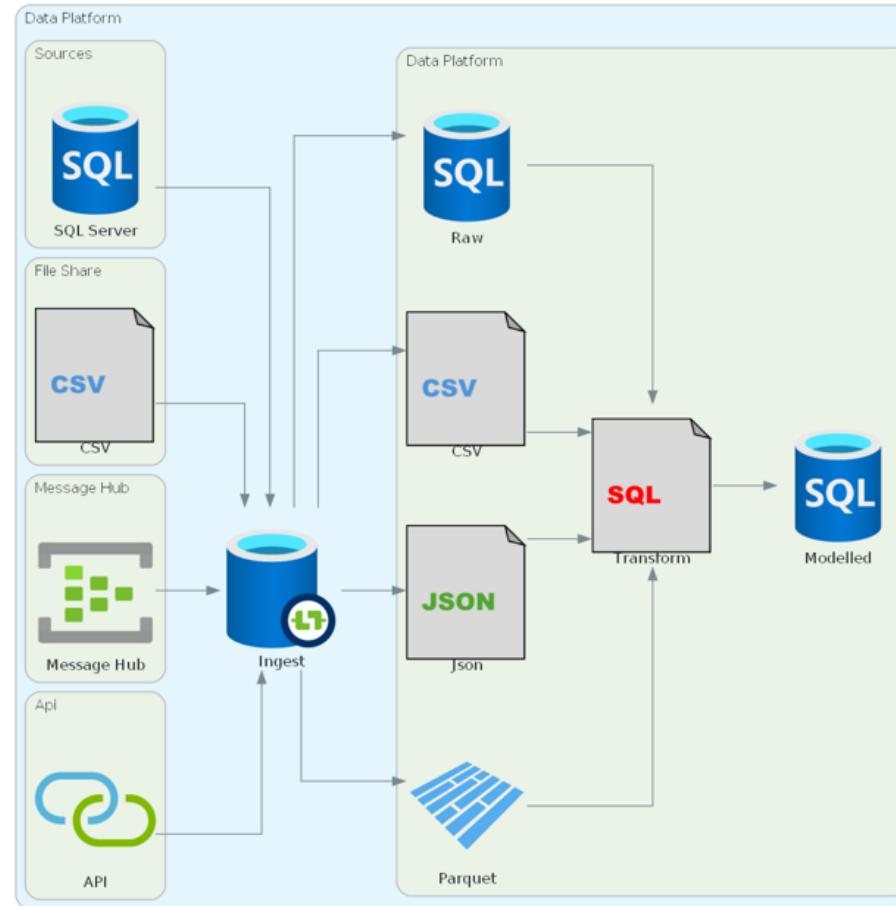
Modern Data Stack

Storage got cheaper



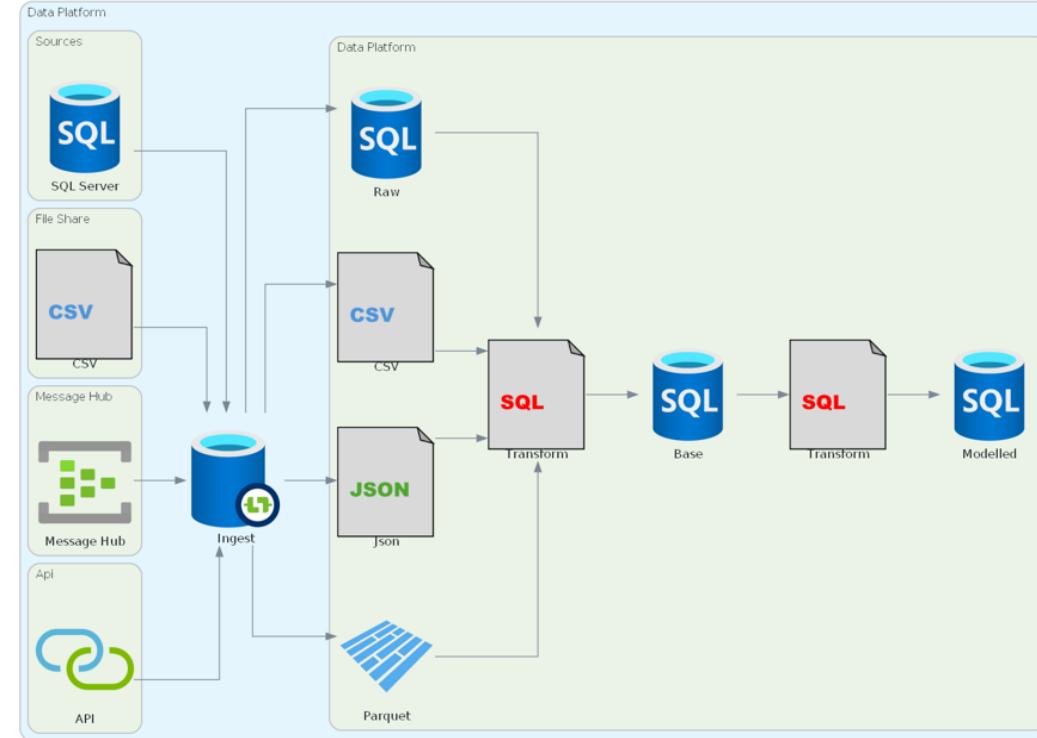
Modern Data Stack

More sources!

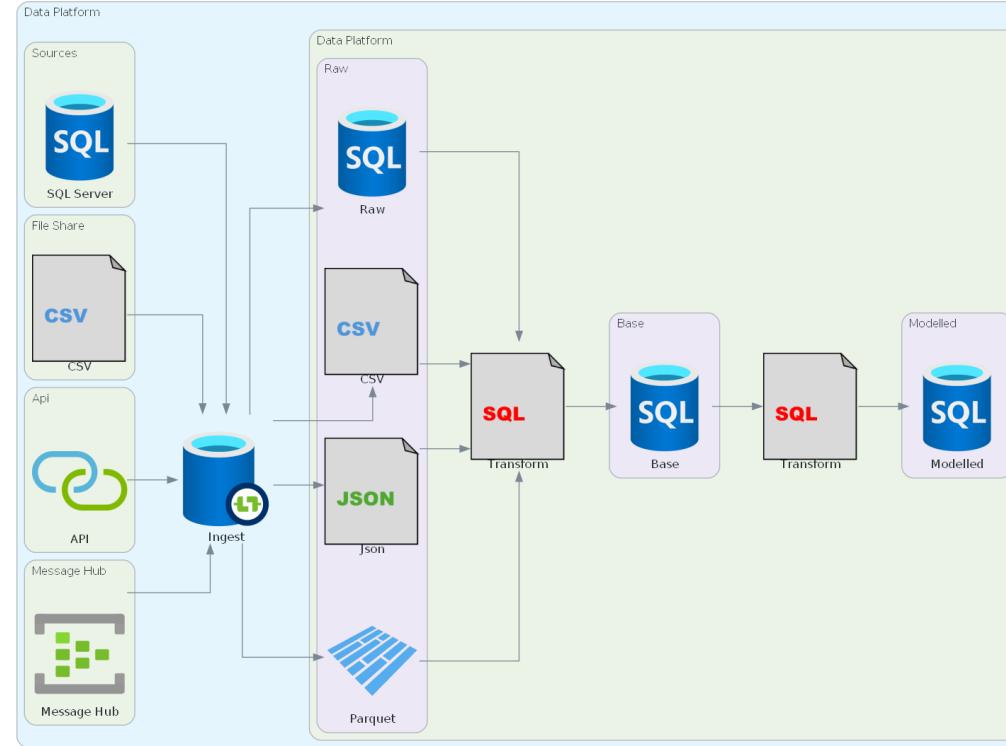


Modern Data Stack

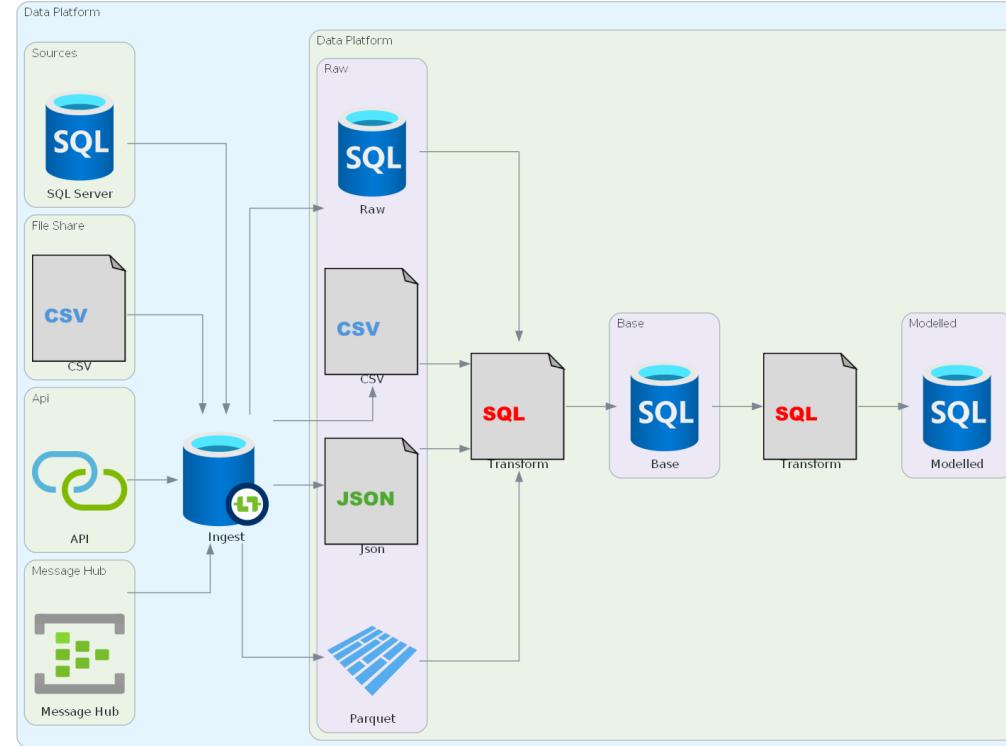
Transformation is hard!



Modern Data Stack

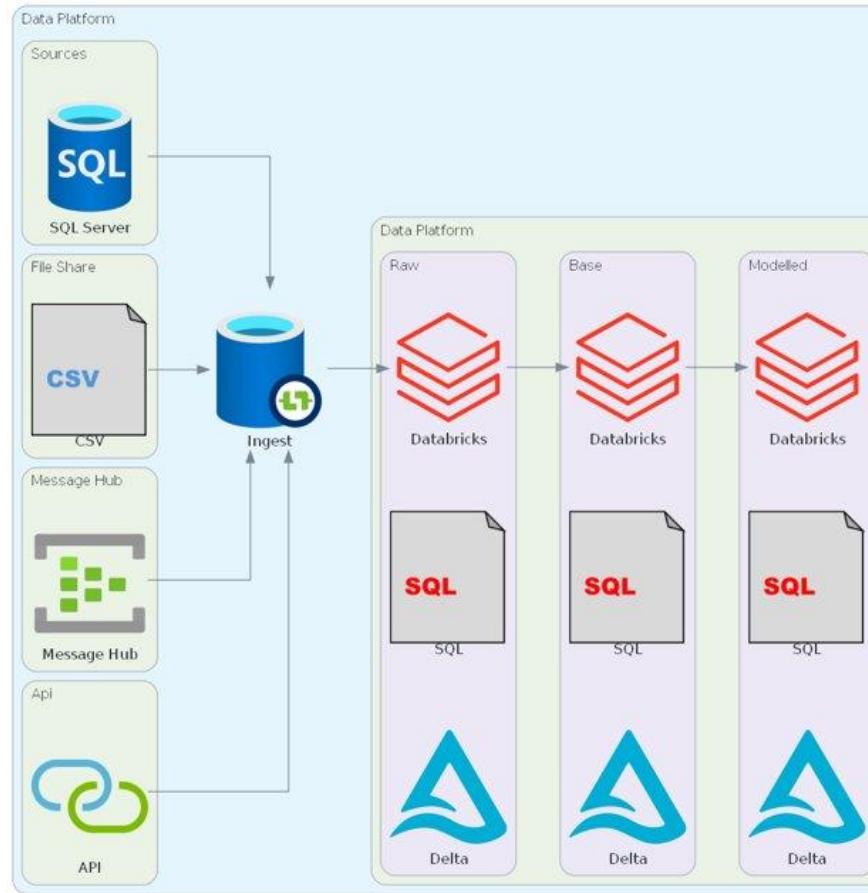


Modern Data Stack



Modern Data Stack

More Saving is on demand compute





Storage & Compute

What problem are we trying to solve?

- What compute engine are we using?

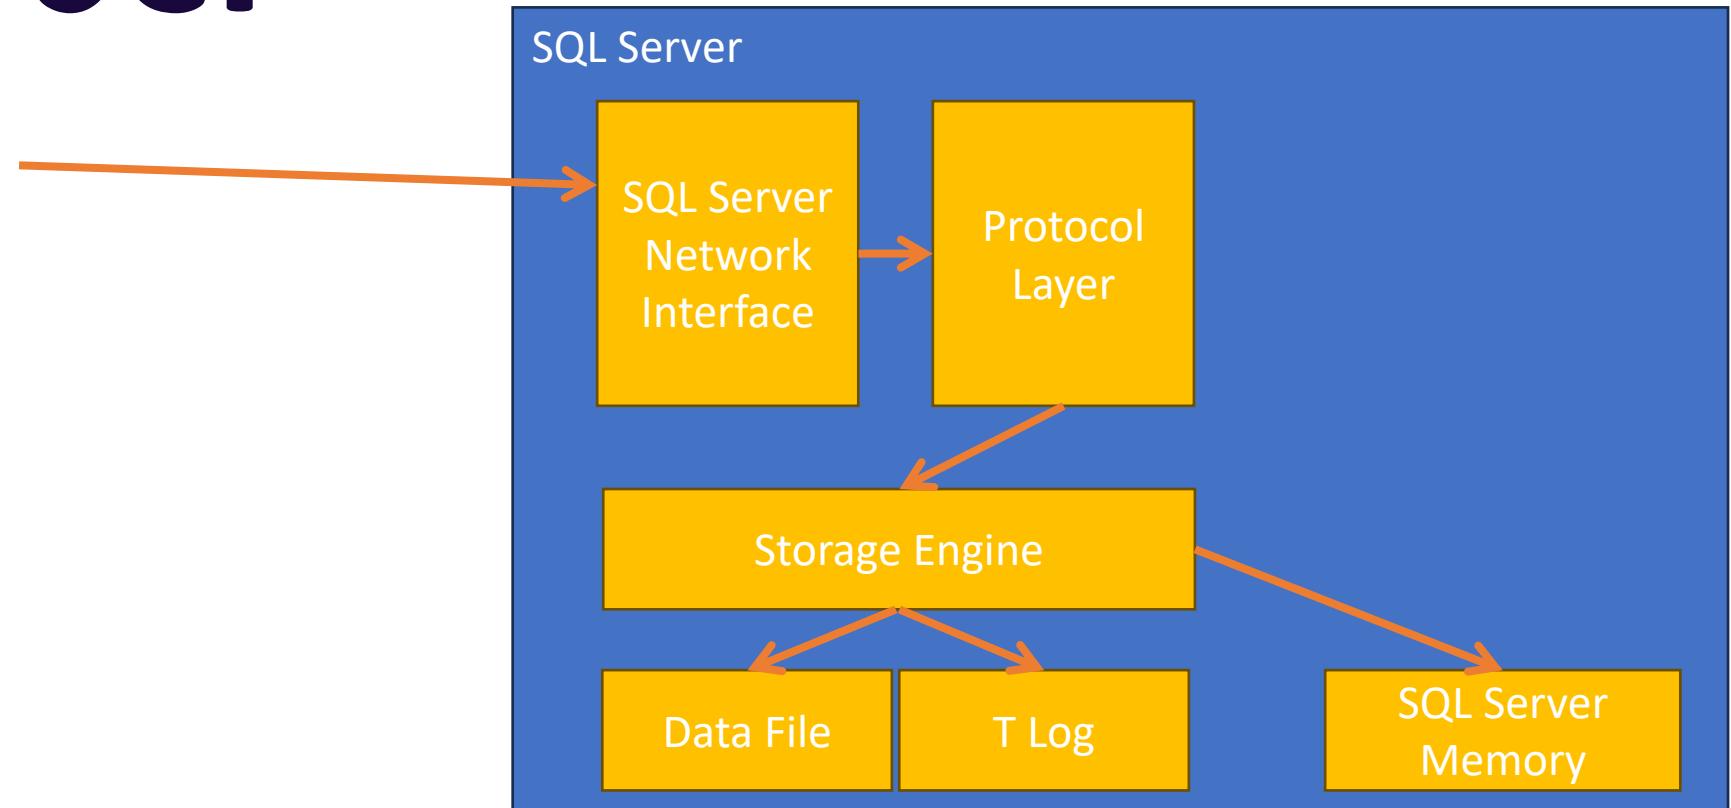


SQL Server

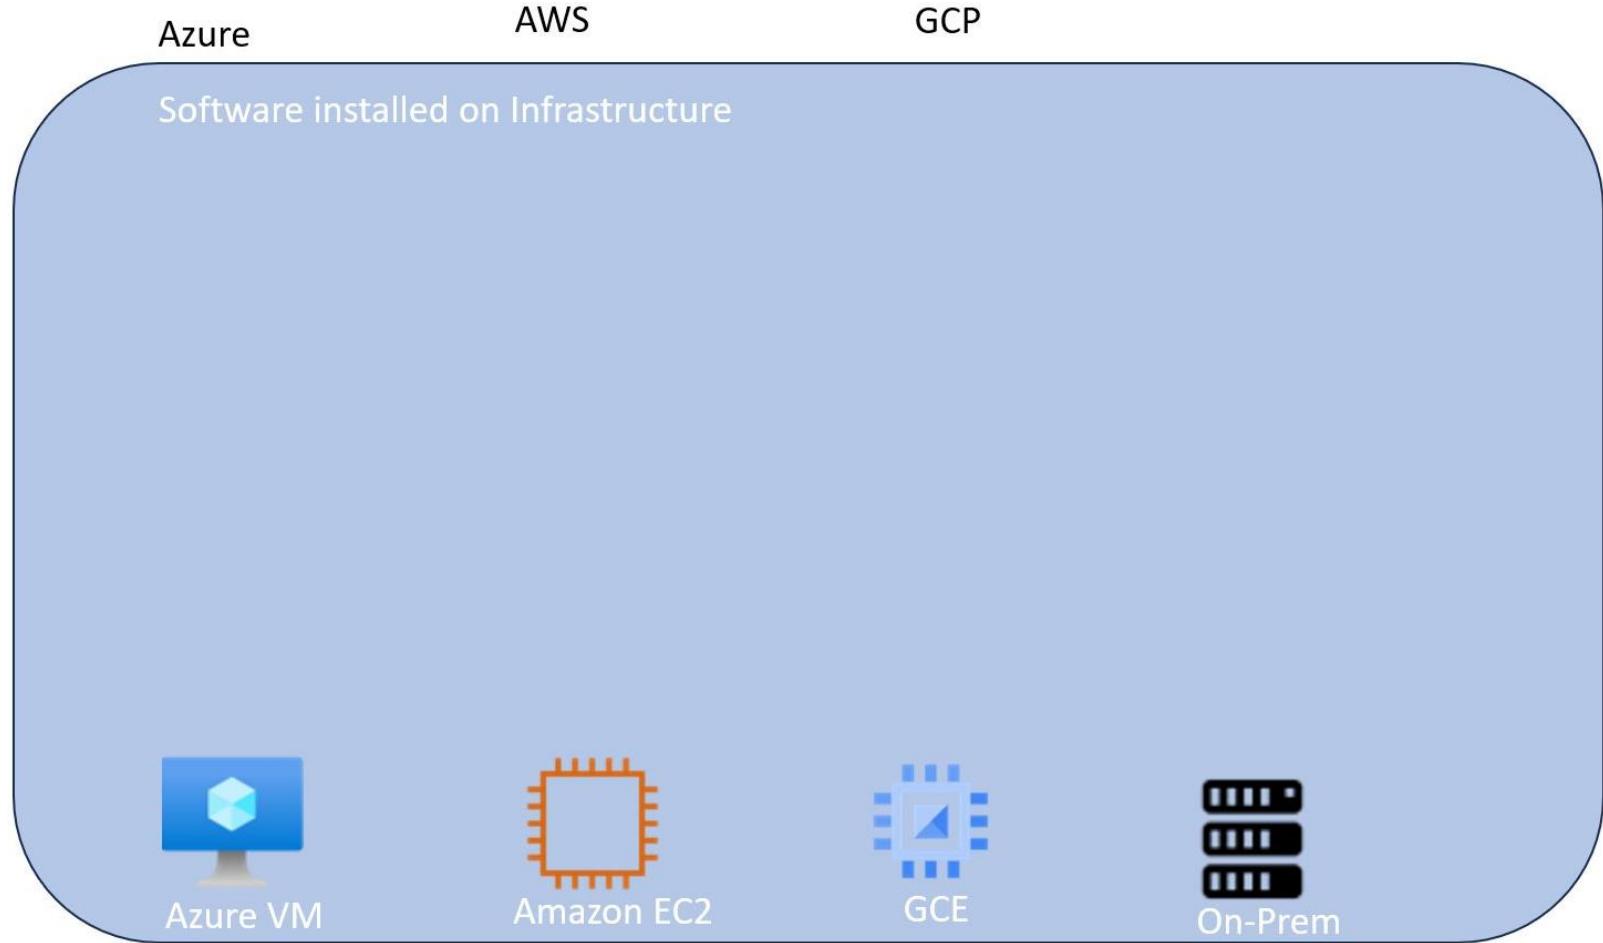
SQL Server

What is Computer?

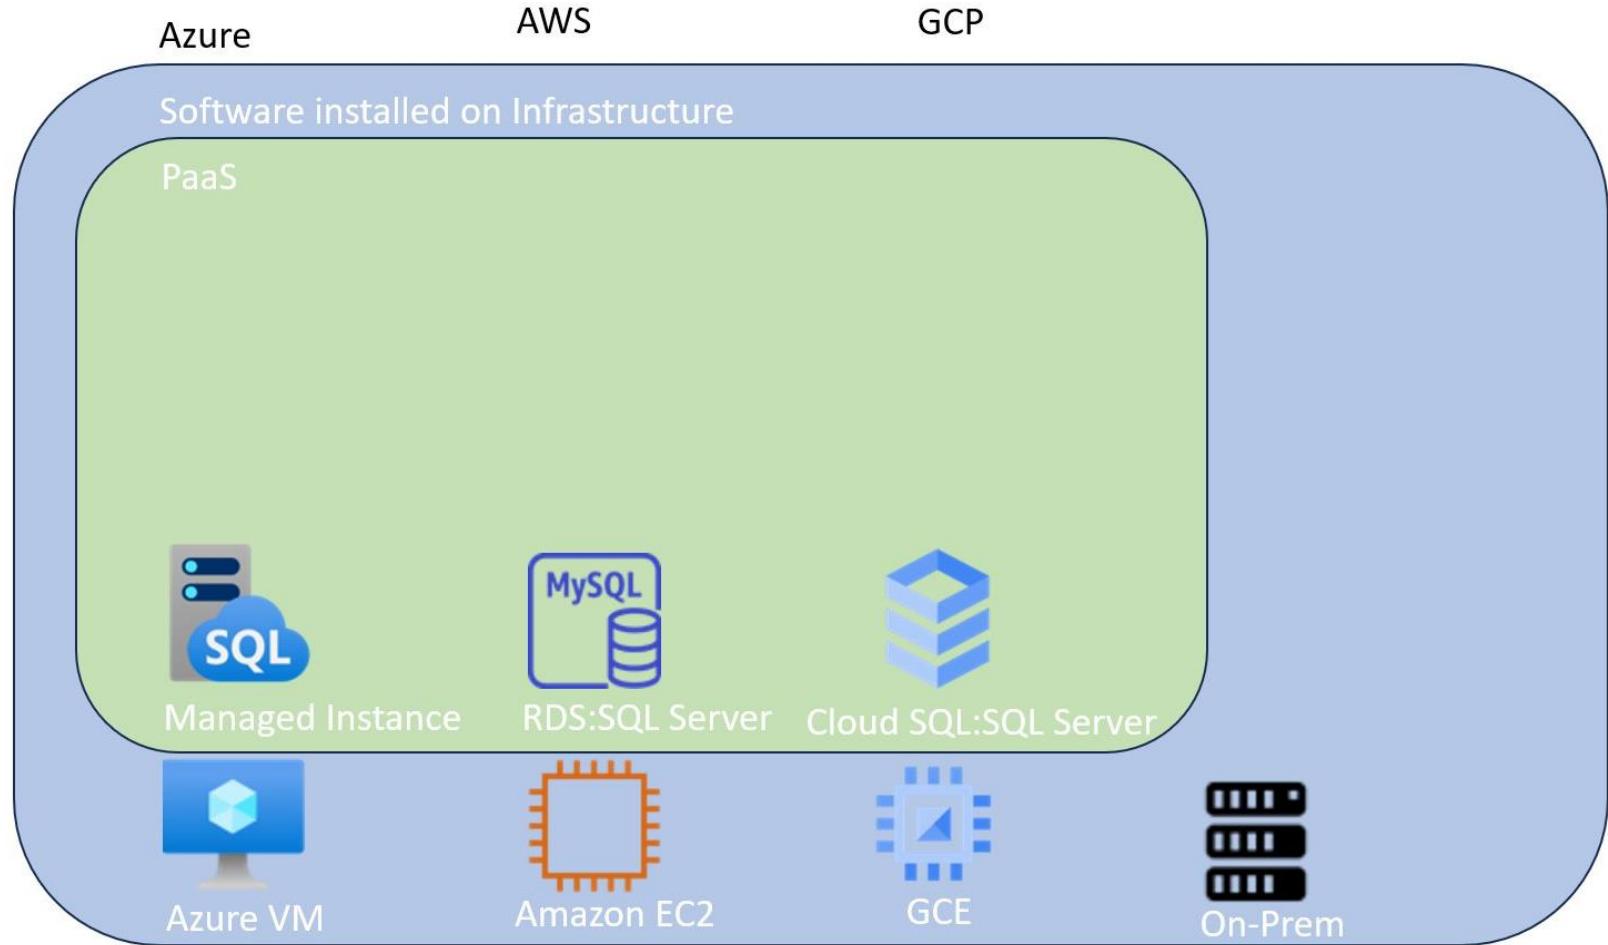
What is Storage?



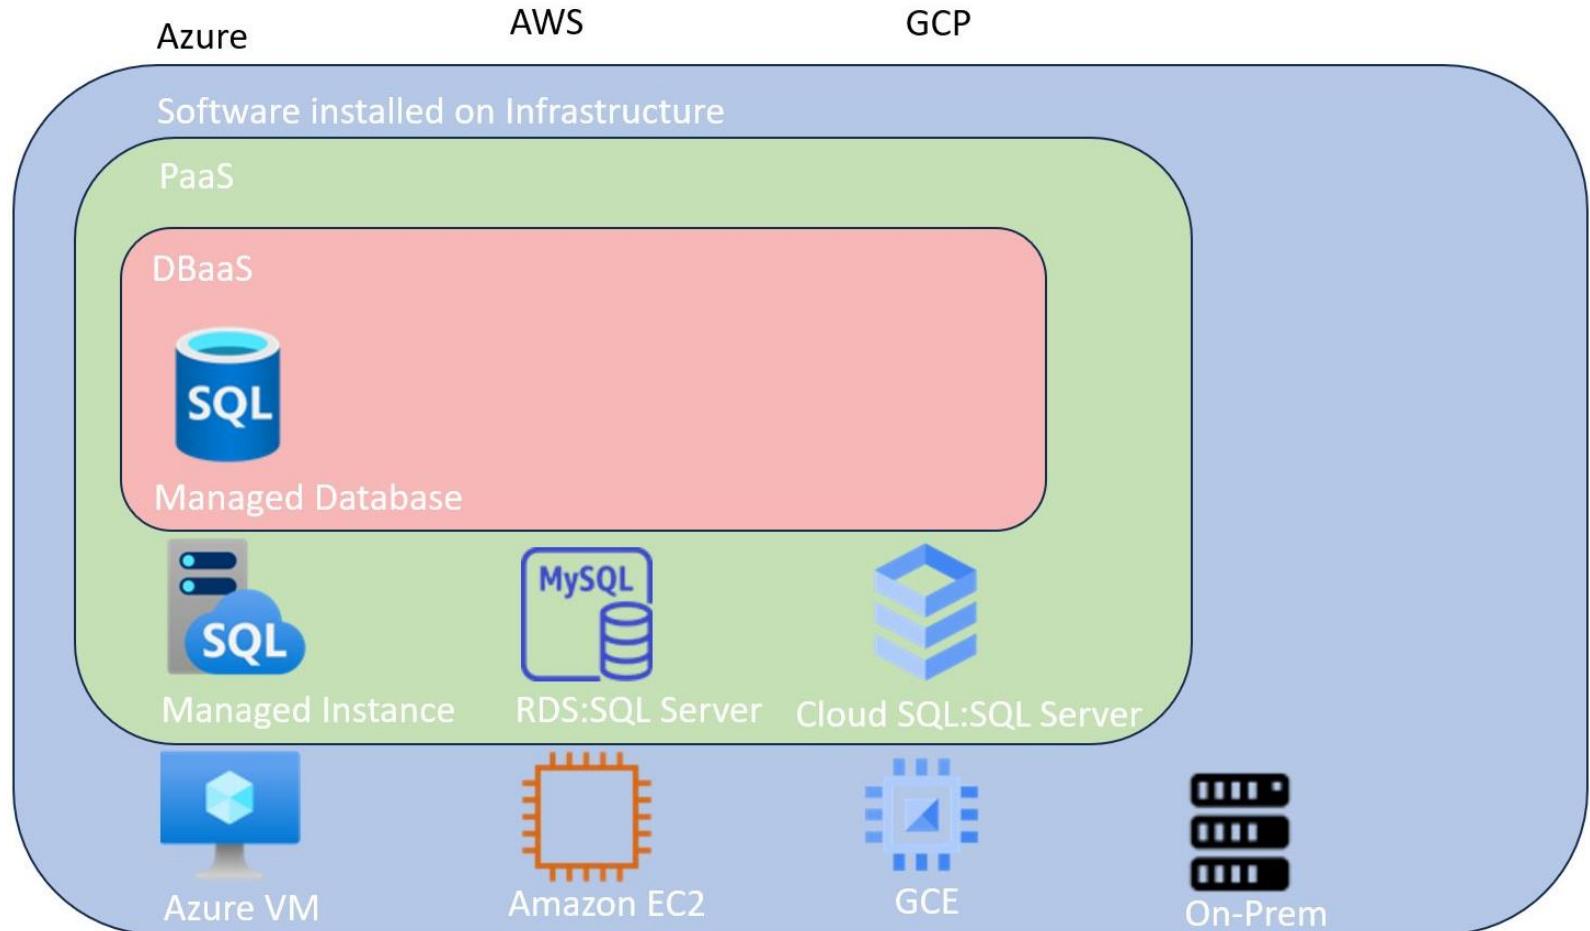
SQL Server



SQL Server



SQL Server



SQL Server

Managed Instance:

- Standard
- Premium
- Memory Optimized Premium

Home > SQL managed instances > Create Azure SQL Managed Instance >

Compute + storage

SQL managed instance

Feedback

Service tier

Select from the latest vCore service tiers available for Azure SQL Managed Instance including General Purpose and Business Critical. [Learn more](#)

Service tier ⓘ

- General Purpose (4-80 vCores, 32 GB-16 TB storage capacity, Fast storage) - for most production workloads
- Business Critical (4-80 vCores, 32 GB-4 TB storage capacity, Super fast storage) - for IO-intensive and compute-intensive workloads

Compute Hardware

Configure compute hardware that will run your Azure SQL Managed Instance. [Learn more](#)

Hardware generation ⓘ

- Standard-series (Gen 5) - Intel Broadwell, 5,1 GB RAM/vCore
- Premium-series - Intel Ice Lake, 7 GB RAM/vCore, up to 560 GB
- Premium-series - memory optimized - Intel Ice Lake, 13,6 GB RAM/vCore, up to 870,4 GB

vCores ⓘ

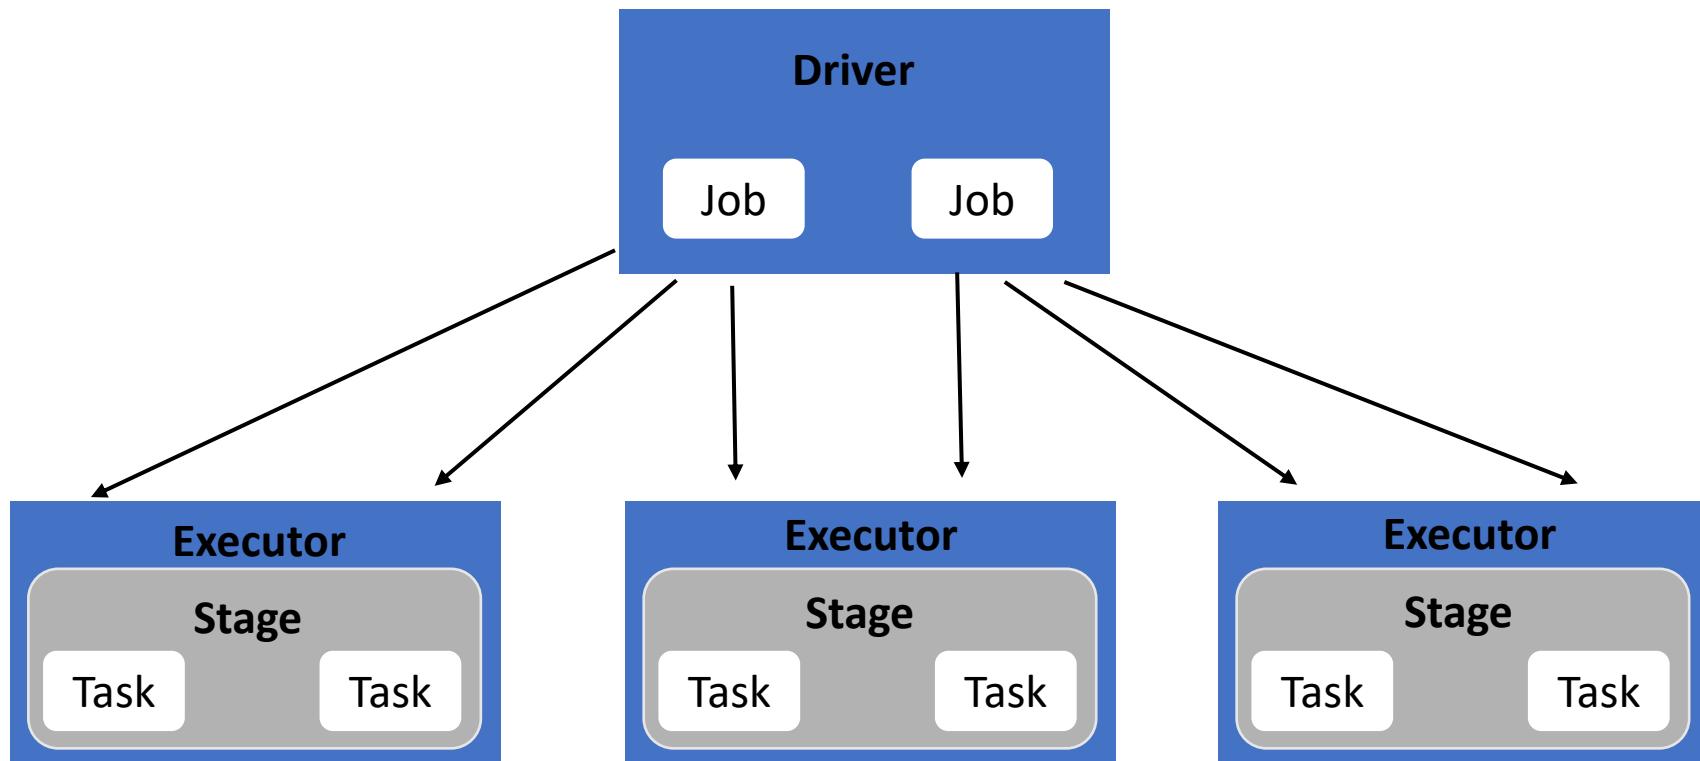


Storage in GB ⓘ



Databricks Clusters

How Databricks Compute Works



Databricks Runtime

Databricks Runtime's are version of **Apache Spark** with **additional libraries** and **performance enhancements** provided by Databricks. Examples include:

- **Databricks Runtime for Apache Spark**
- **Databricks Runtime ML**
- **Databricks Runtime for SQL Analytics**

Different types of Databricks Compute

- All-Purpose compute
- Job compute
- Instance pools
- Serverless SQL warehouses



Create a Databricks Cluster

Demo Cluster [🔗](#)

Performance

Databricks runtime version [?](#)

Runtime: 13.3 LTS (Scala 2.12, Spark 3.4.1) [|](#) [▼](#)

Use Photon Acceleration [?](#)

Worker type [?](#)

Standard_DS3_v2 14 GB Memory, 4 Cores [|](#) [▼](#)

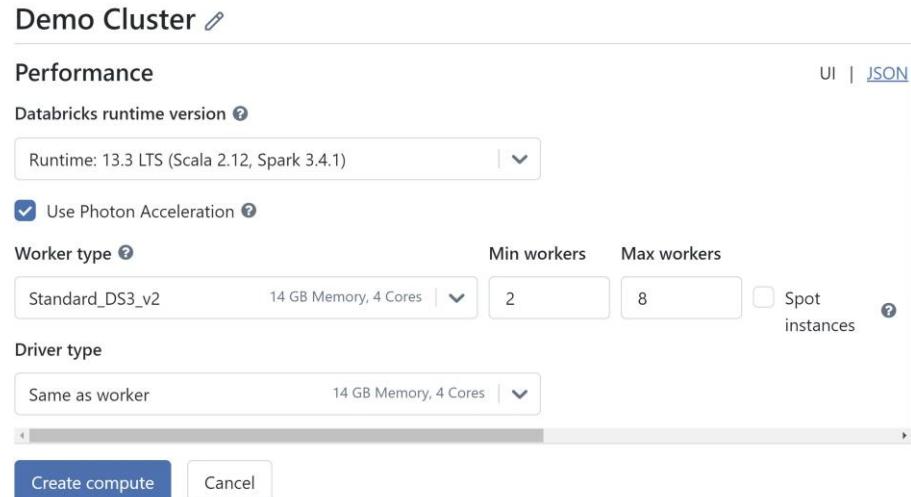
Min workers: 2 Max workers: 8 Spot instances [?](#)

Driver type

Same as worker 14 GB Memory, 4 Cores [|](#) [▼](#)

[Create compute](#) [Cancel](#)

UI | [JSON](#) [?](#)



Databricks supports creating and managing clusters with different configurations based on the **cluster mode**, **instance type**, and **autoscaling**. Standard clusters can be used for ad-hoc data exploration, batch jobs, and data engineering tasks.

Photon

Databricks-native vectorized query engine that runs your **SQL workloads** and **DataFrame API calls** faster



Limitations

- Does **NOT support** Structured Streaming: Photon currently supports **stateless streaming ONLY** with Delta, Parquet, CSV, and JSON. Stateless Kafka and Kinesis streaming is supported when writing to a Delta or Parquet sink.
- Does **NOT support UDFs or RDD APIs**.
- **WONT** impact queries that normally run in under two seconds

ACID Transactions

Atomicity, Consistency, Isolation, and Durability

A
tomicity

The Transaction can be treated as a single unit
that can succeed or fail

C

I

D



Atomicity, Consistency, Isolation, and Durability

A Atomicity The Transaction can be treated as a single unit
 that can succeed or fail

C Consistency The data must be consistent before and
 after the Transaction

I

D



Atomicity, Consistency, Isolation, and Durability

Atomicity The Transaction can be treated as a single unit that can succeed or fail

Consistency The data must be consistent before and after the Transaction

Isolation The Transaction is independent and can be attempted without interference

D



Atomicity, Consistency, Isolation, and Durability

Atomicity	The Transaction can be treated as a single unit that can succeed or fail
Consistency	The data must be consistent before and after the Transaction
Isolation	The Transaction is independent and can be attempted without interference
Durability	A successful transaction persists if the system fails

Atomicity, Consistency, Isolation, and Durability

	Atomicity	Consistency	Isolation	Durability
Parquet	✗	✗	✗	✗
Delta Format	✓	✓	✓	✓
Simple Recovery	✓	✓	✓	✓
Full Recovery	✓	✓	✓	✓

Isolation Levels

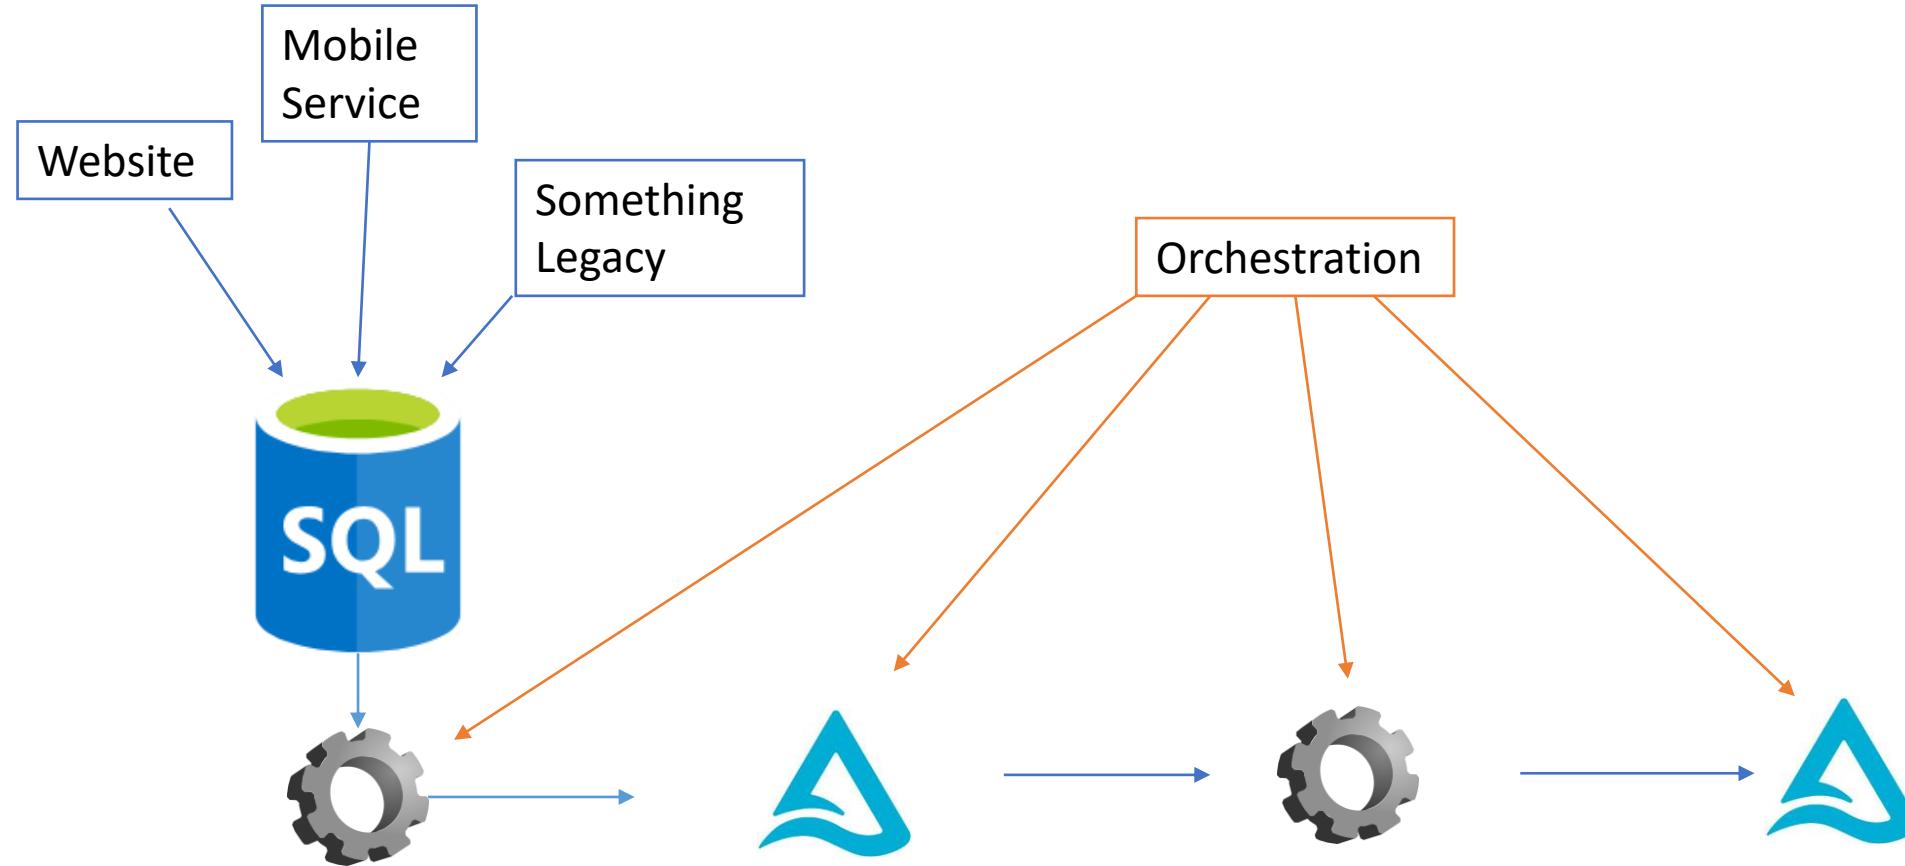
SQL Server	Delta Format
Read Uncommitted	Write Serializable
Repeatable Read	✗
Snapshot	✗
Serializable	Serializable

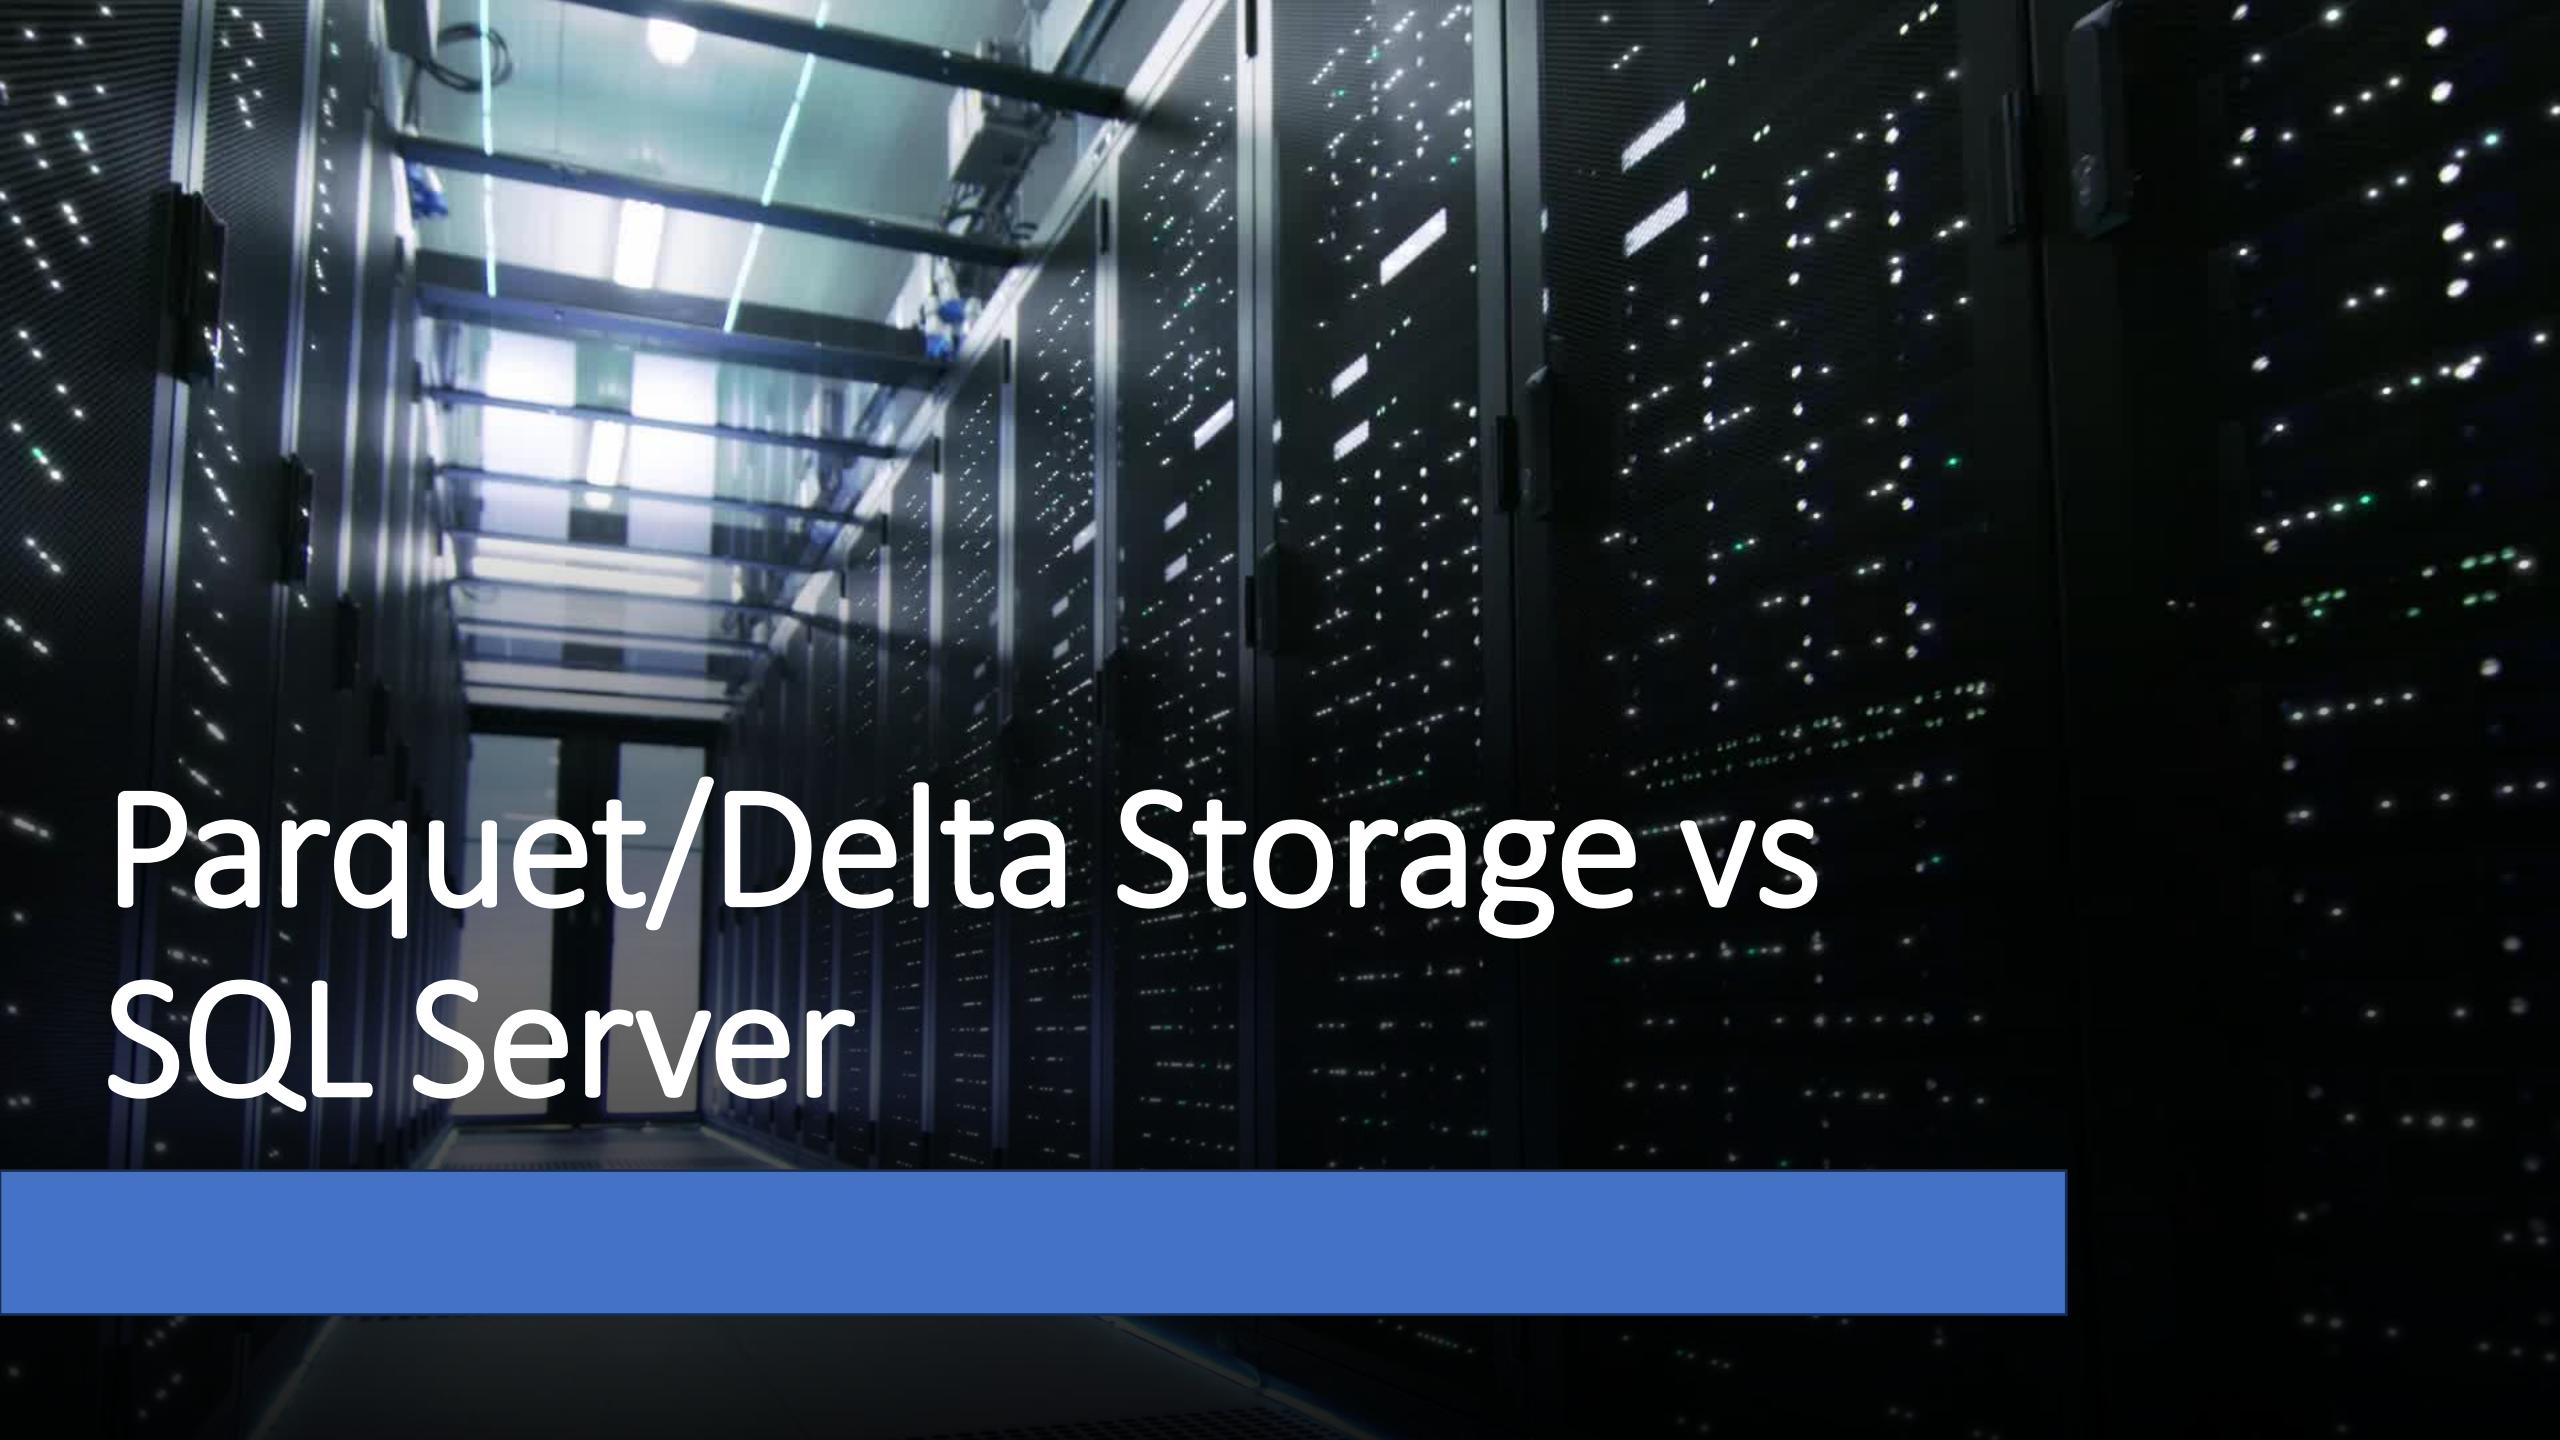
Atomicity, Consistency, Isolation, and Durability

- Databricks uses Delta format to get ACID properties
- Delta format is open source!



Modern Data Stack





Parquet/Delta Storage vs SQL Server

What problem are we trying to solve?

Where and how to store our data, what format do we use?



Storage in SQL Server

How do we store data to disk?

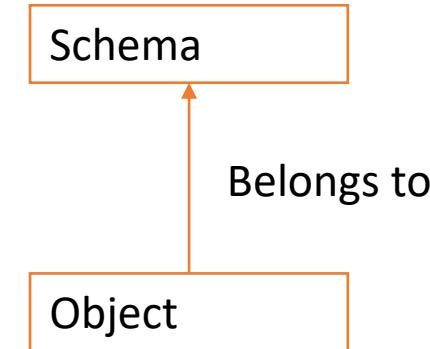
We have to persist data to disk

- Data stored in Data files
- SQL Server exclusively access the files

	SQL Server	Delta Format	Parquet
File Structure	*.MDF, *.NDF	*.parquet	*.parquet
Log File Structure	*.LDF	/_delta_log /000000.json	

How do we store data to disk?

```
CREATE TABLE [Person].[Person]([BusinessEntityID] [int] NOT NULL,  
[PersonType] [nchar](2) NOT NULL,  
[NameStyle] [nvarchar](15) NOT NULL,  
[Title] [nvarchar](8) NULL,  
[FirstName] [nvarchar](50) NOT NULL,  
[MiddleName] [nvarchar](50) NULL,  
[LastName] [nvarchar](50) NOT NULL,  
[Suffix] [nvarchar](10) NULL,  
[EmailPromotion] [int] NOT NULL,  
[AdditionalContactInfo] [xml] NULL,  
[Demographics] [xml] NULL,  
[rowguid] [uniqueidentifier] ROWGUIDCOL NOT NULL,  
[ModifiedDate] [datetime] NOT NULL,  
) ON [PRIMARY] GO
```



How do we store data to disk?

Object created on Primary will be in
the primary file group

Database Object

Database File

sp_spaceused 'TableName'

[Person].[Person]

MDF

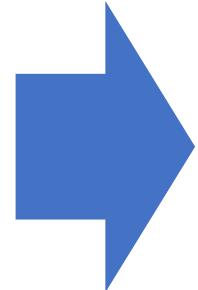
```
SELECT *  
FROM sys.filegroups;
```

Parquet/Delta Storage

What is Parquet?

Open source, column-oriented data **file format**

A	1	Fred
A	1	Bob
A	1	Fred
B	1	Bob
C	1	Fred
C	1	Bob
C	1	Fred
B	1	Bob



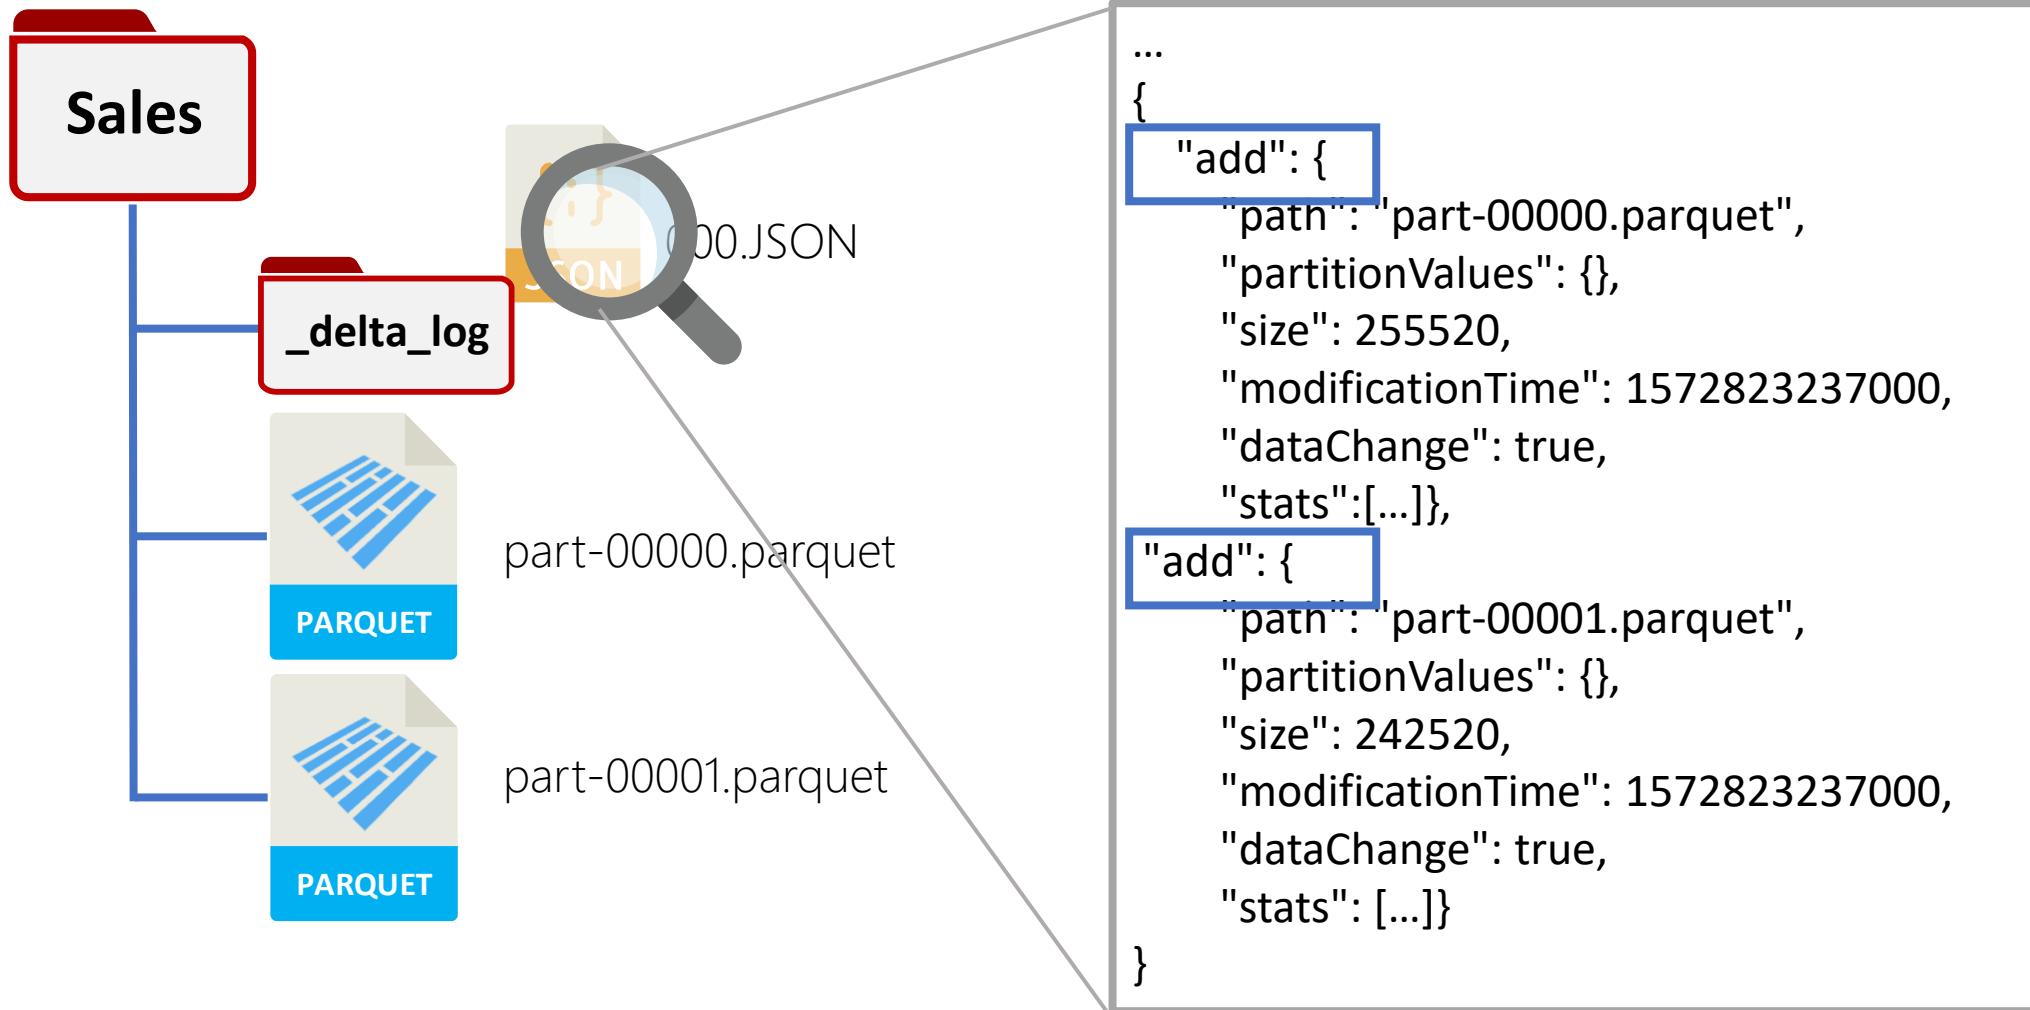
CSV File

Col1	Col2	Col3
A*3	1*8	1
B		2
C*3		1
B		2
		1
		2
		1
		2

Parquet File

Parquet works in a **very similar way to SQL ColumnStore**

It's a **ColumnStore compression format**



Delta Features

- ACID Transactions
- Scalable Metadata
- Time Travelling
- Audit History
- Batch & Streaming Support
- Schema Enforcement
- Schema Evolution
- Familiar SQL Commands
- Open Format
- Compatible with Spar



SQL Storage Structures

Storage Structures

Within these files, SQL Server has three types of storage structures that are Heap, clustered and column store indexes

- Heap
- Clustered Index
- Column store Index

This are where the data is physically stored on Disk

Storage Structures

We can mix these to meet difference performance benefits

Let's stay with person.person

We have different access patterns from different queries

Let's try and search for some by [BusinessEntityID]

Storage Structures

```
SELECT [BusinessEntityID]
      ,[PersonType]
      ,[NameStyle]
      ,[Title]
      ,[FirstName]
      ,[MiddleName]
      ,[LastName]
      ,[Suffix]
      ,[EmailPromotion]
      ,[AdditionalContactInfo]
      ,[Demographics]
      ,[rowguid]
      ,[ModifiedDate]
  FROM [AdventureWorks2022].[Person].[Person]
 WHERE [BusinessEntityID] = 4000
```

SQL Server Engine

[PK_Person_BusinessEntityID]

[IX_Person_LastName_FirstName_MiddleName]

[AK_Person_rowguid]

[PXML_Person_AddContact]

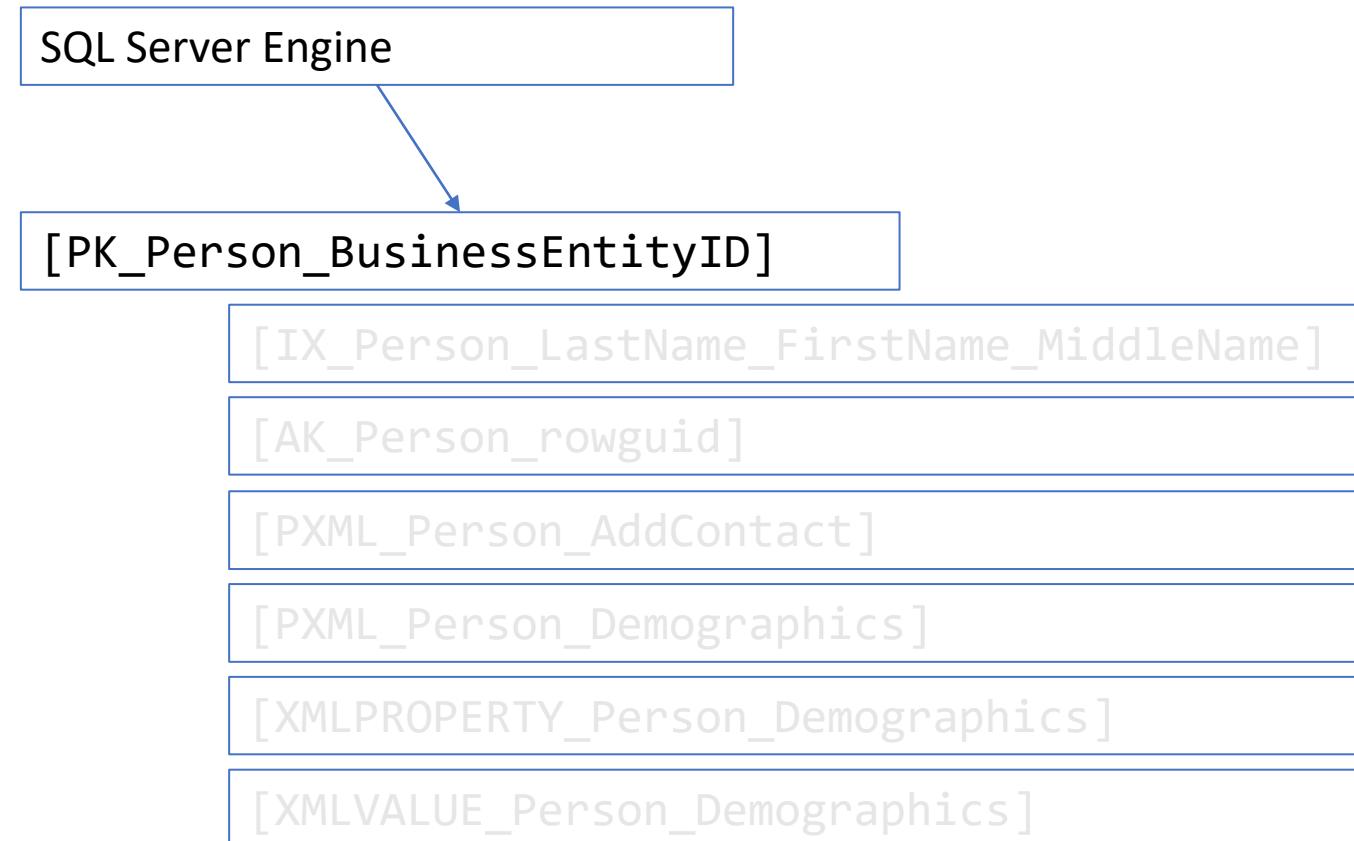
[PXML_Person_Demographics]

[XMLPROPERTY_Person_Demographics]

[XMLVALUE_Person_Demographics]

Storage Structures

```
SELECT [BusinessEntityID]
      ,[PersonType]
      ,[NameStyle]
      ,[Title]
      ,[FirstName]
      ,[MiddleName]
      ,[LastName]
      ,[Suffix]
      ,[EmailPromotion]
      ,[AdditionalContactInfo]
      ,[Demographics]
      ,[rowguid]
      ,[ModifiedDate]
  FROM [AdventureWorks2022].[Person].[Person]
 WHERE [BusinessEntityID] = 4000
```



Storage Structures

SQL Server Engine

```
SELECT [BusinessEntityID]
    ,[PersonType]
    ,[NameStyle]
    ,[Title]
    ,[FirstName]
    ,[MiddleName]
    ,[LastName]
    ,[Suffix]
    ,[EmailPromotion]
    ,[AdditionalContactInfo]
    ,[Demographics]
    ,[rowguid]
    ,[ModifiedDate]
FROM [AdventureWorks2022].[Person].[Person]
WHERE [FirstName] = 'Jeremiah'
AND [LastName] = 'Martin'
```

[PK_Person_BusinessEntityID]

[IX_Person_LastName_FirstName_MiddleName]

[AK_Person_rowguid]

[PXML_Person_AddContact]

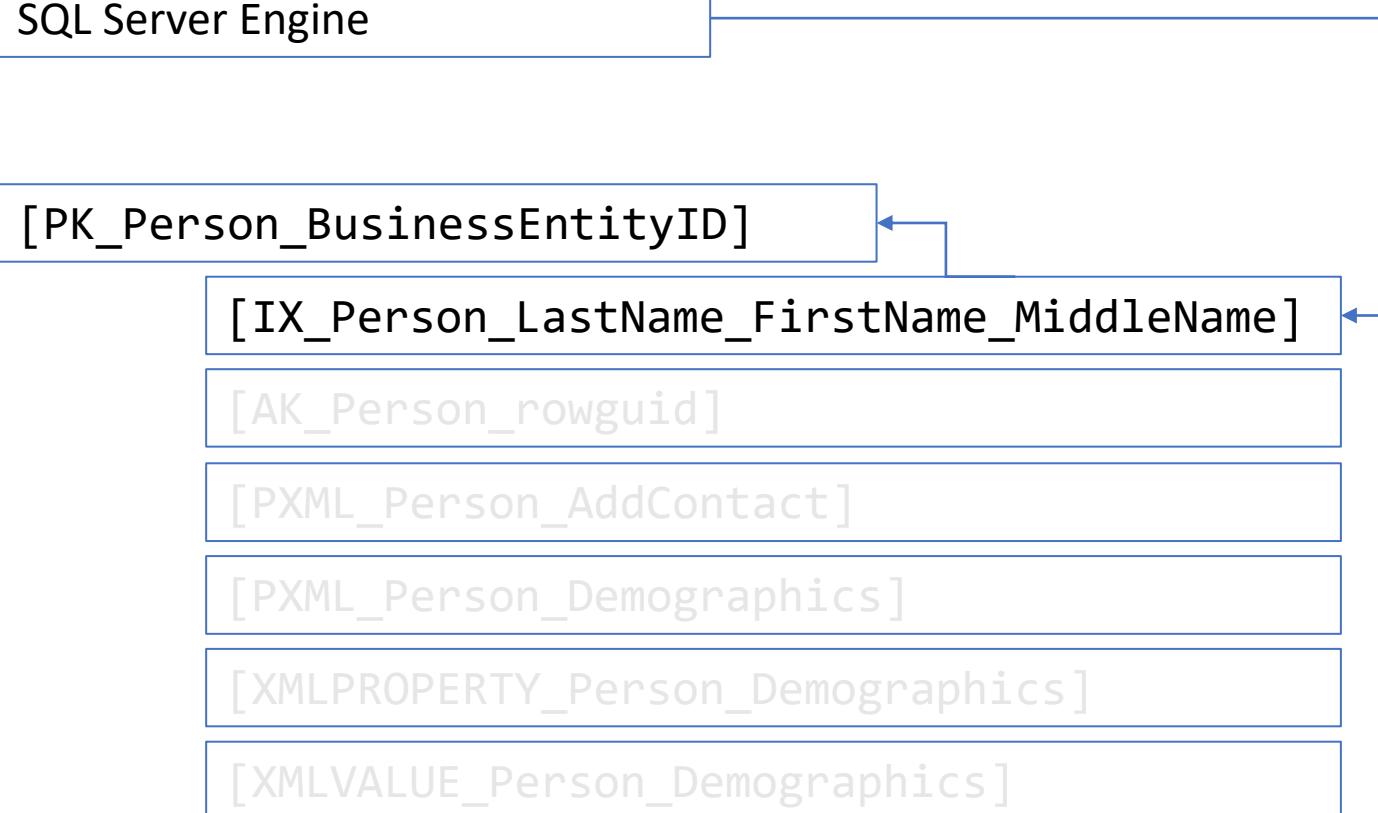
[PXML_Person_Demographics]

[PROPERTY_Person_Demographics]

[VALUE_Person_Demographics]

Storage Structures

```
SELECT [BusinessEntityID]
      ,[PersonType]
      ,[NameStyle]
      ,[Title]
      ,[FirstName]
      ,[MiddleName]
      ,[LastName]
      ,[Suffix]
      ,[EmailPromotion]
      ,[AdditionalContactInfo]
      ,[Demographics]
      ,[rowguid]
      ,[ModifiedDate]
  FROM [AdventureWorks2022].[Person].[Person]
 WHERE [FirstName] = 'Jeremiah'
   AND [LastName] = 'Martin'
```



Storage Structures

SQL Server allow us to have a transactional copy of the data as a non-clustered index that can optimise reads.

In contrast, we would create copy with a subset in Databricks!

	SQL Server	Delta	Parquet
Data Structure	Heap, B Tree, Clustered Column Store, Order Clustered Column Store	Columnal	Columnal
Secondary Data Copies	Non Clustered Index, Non Clustered Column Store		



Storage Structures

We can have multiple, transactionally identical structure to meet our use cases

We can index for row level and aggregate queries allow SQL Server to meet this use cases individually or concurrently!

But this does have its own overheads!





Partitioning in Delta and SQL Server

What problem are we solving?

- Chunking data into logical buckets for reading and writing without having to work with the entire dataset
- Large Data Set can have issues at scale with data manipulation and storage size



Partitioning in SQL Server

Partition View and Tables

SQL Server has two type of partitioning.

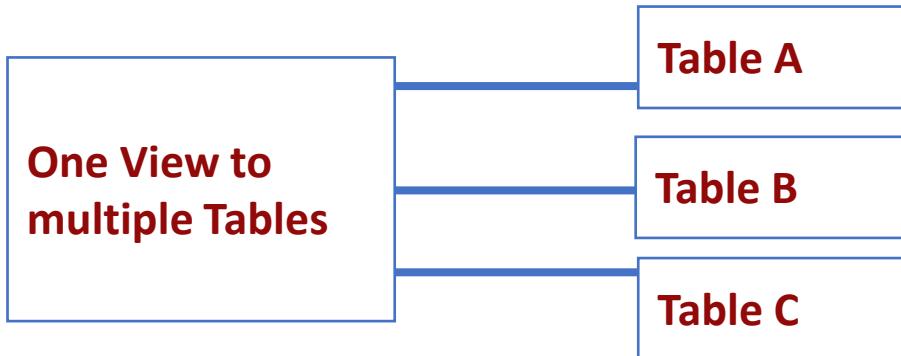
- Partition Table
- Partition Views

Partition Tables are tables partitioned on a key across multiple files.

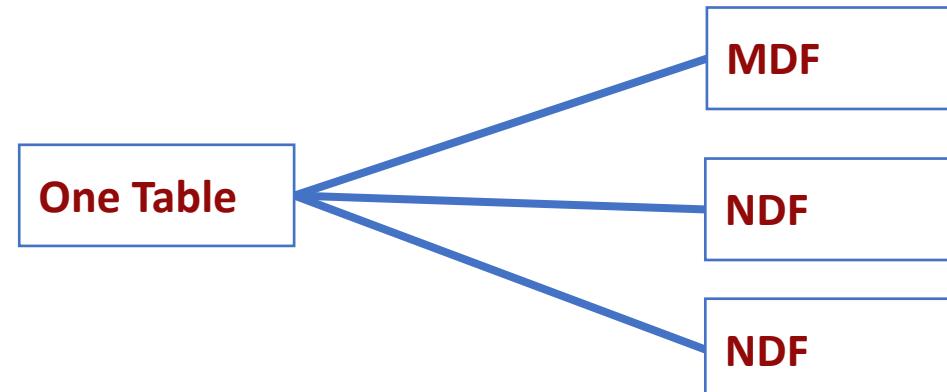
Partition views can be place over tables to partition different keys.

Partition View and Tables

V
i
e
w

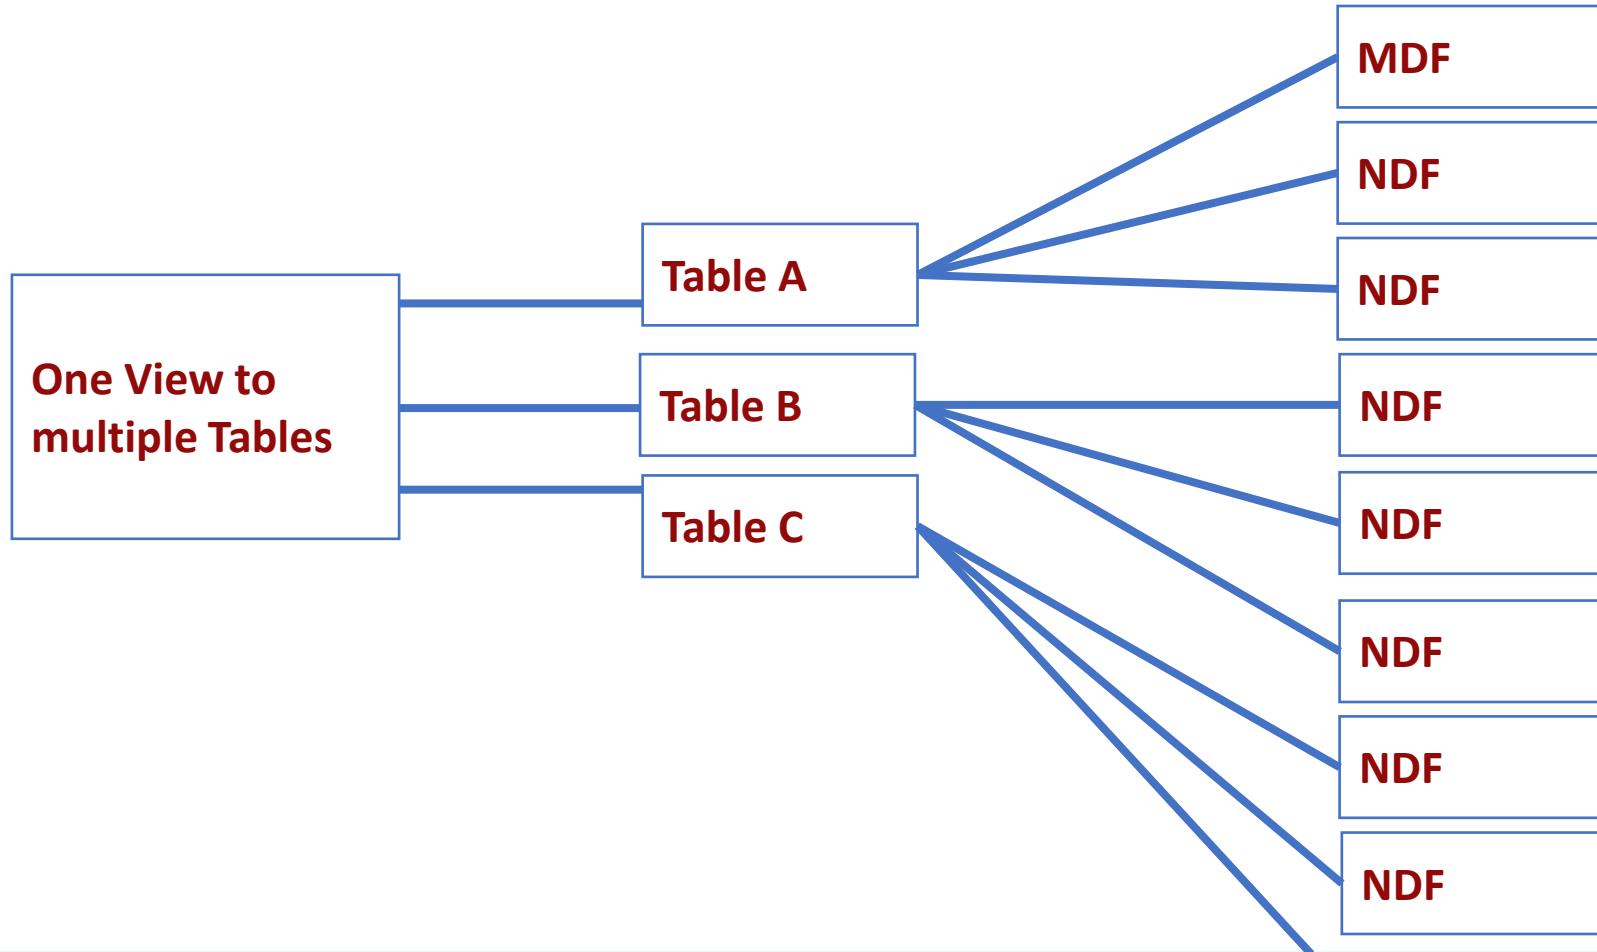


T
a
b
l
e



Each table will have its own file group

Partition View and Tables



Partition View and Tables

So, lets look at adventure works, let looks at the [Purchasing].[PurchaseOrderHeader] table

```
CREATE TABLE [Purchasing].[PurchaseOrderHeader](
[PurchaseOrderID] [int] IDENTITY(1,1) NOT NULL,
[RevisionNumber] [tinyint] NOT NULL,
[Status] [tinyint] NOT NULL,
```

Partition View and Tables

So, lets look at adventure works, let looks at the [Purchasing].[PurchaseOrderHeader] table

```
CREATE TABLE [Purchasing].[PurchaseOrderHeader_1](
[PurchaseOrderID] [int] INTEGER DEFAULT NEXT VALUE FOR PurchaseOrderHeader_Seq,
[RevisionNumber] [tinyint] NOT NULL,
[Status] [tinyint] NOT NULL CONSTRAINT [PK_PurchaseOrderHeader_Status_1] CHECK ( [Status]
= 1 ),
```

Partition View and Tables

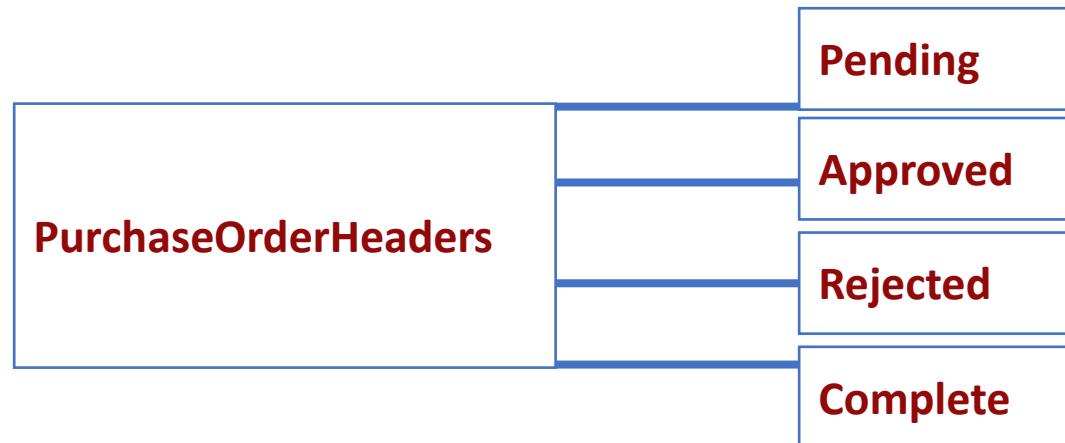
We can start with a partition view, which will be a view that spans the tables

We already have [PurchaseOrderID]

We can now also use the [Status] so that we can split out, Pending, Approved, Rejected and Complete

Partition View and Tables

So we will now create these tables with additional constraints so that the data will be in the correct table



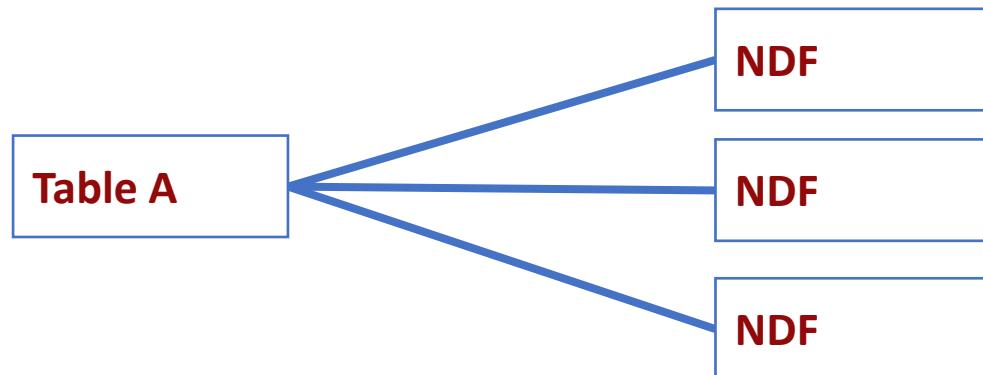
Partition View and Tables

Demo Time!

Partition View and Tables

We now have a partition view made of tables, this tables all have the same schema and so we will apply partitioning to them in a group.

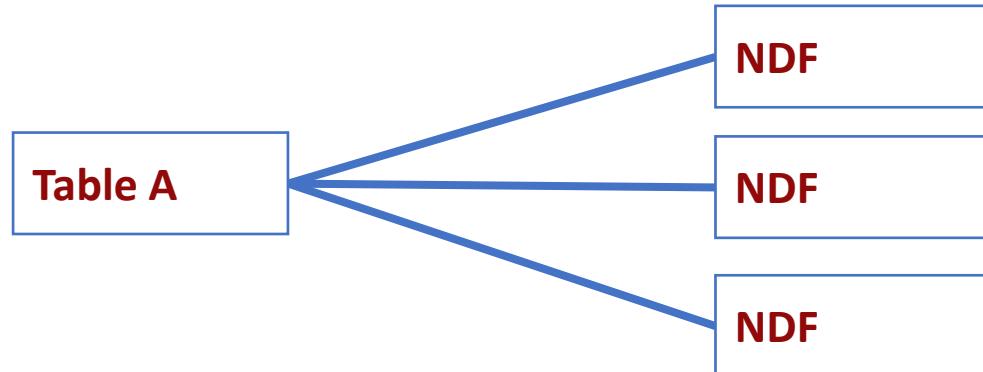
We need to create a Partition Schema and Partition Function



Partition View and Tables

Our partition Schema is where things can be put!

Our Partition Function how what is going to be based on!



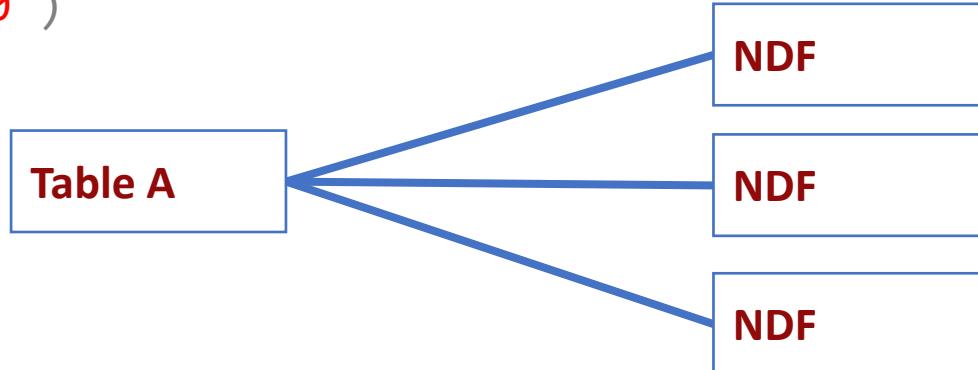
Partition View and Tables

```
CREATE PARTITION SCHEME [PartSch_Purchasing_PurchaseOrderHeader] AS PARTITION  
[PartFun_Purchasing_PurchaseOrderHeader] TO ([Part_2011], [Part_2012], [Part_2013],  
[Part_2014])
```

GO

```
CREATE PARTITION FUNCTION [PartFun_Purchasing_PurchaseOrderHeader](datetime) AS RANGE  
RIGHT FOR VALUES (N'2012-01-01T00:00:00.000', N'2013-01-01T00:00:00.000', N'2014-01-  
01T00:00:00.000')
```

GO



Partition View and Tables

Demo Time!

File Edit View Help

CONNECTIONS ...

Servers <default> (Inte...

C:\> Users > James > SqlBits-2024-Demo1.ipynb

+ Cell ▾ Run all Kernel SQL Attach to AdventurePlay (Integrated) ▾

Create Schema and Sequence

```
[11] 1 IF NOT EXISTS (SELECT * FROM sys.schemas WHERE name = 'Purchasing')
2 BEGIN
3 EXEC('CREATE SCHEMA [Purchasing]')
4 END
5 /*Create Seq*/
6 IF EXISTS (SELECT * FROM sys.sequences WHERE name = 'Purchasing')
7 BEGIN
8 CREATE SEQUENCE PurchaseOrderHeader_Seq
9 AS INTEGER
10 START WITH 1
11 INCREMENT BY 1
12 END
13 /*Drop view fro schema binding*/
14 IF EXISTS (SELECT * FROM sys.views WHERE name = 'PurchaseOrderHeader')
15 BEGIN
16 EXEC('DROP VIEW [Purchasing].[PurchaseOrderHeader]')
17 END
```

Commands completed successfully.

Total execution time: 00:00:00.035

Create New tables for Partitioning across status

All tables on Primary as this is part of a Demo

```
[12] 1 /*Demo Starts here*/
2 /* Create Table*/
3 IF EXISTS (SELECT * FROM sys.tables WHERE name = 'PurchaseOrderHeader_1')
4 BEGIN
5 EXEC('DROP TABLE [Purchasing].[PurchaseOrderHeader_1]')
6 END
7
8 CREATE TABLE [Purchasing].[PurchaseOrderHeader_1](
9 [PurchaseOrderID] INTEGER DEFAULT NEXT VALUE FOR PurchaseOrderHeader_Seq,
10 [RevisionNumber] [tinyint] NOT NULL,
11 [Status] [tinyint] NOT NULL CONSTRAINT [PK_PurchaseOrderHeader_Status_1] CHECK ( [Status] = 1 ),
12 [EmployeeID] [int] NOT NULL,
13 [VendorID] [int] NOT NULL,
14 [ShipMethodID] [int] NOT NULL,
15 [OrderDate] [datetime] NOT NULL,
16 [ShipDate] [datetime] NULL,
17 [SubTotal] [money] NOT NULL,
18 [TaxAmt] [money] NOT NULL,
19 [Freight] [money] NOT NULL,
20 --[TotalDue] AS (ISNULL((SubTotal)+(TaxAmt)),(0)) PERSISTED NOT NULL,
21 [ModifiedDate] [datetime] NOT NULL,
22 CONSTRAINT [PK_PurchaseOrderHeader_PurchaseOrderID_1] PRIMARY KEY CLUSTERED
23 (
24 [PurchaseOrderID],[Status] ASC
25 ) ON [PRIMARY]
26 ) ON [PRIMARY]
27 GO
```

Commands completed successfully.

Total execution time: 00:00:00.049

Repeate for PurchaseOrderHeader_2, PurchaseOrderHeader_3 and PurchaseOrderHeader_4

SQL

SQL



Partitioning in Delta

Delta Partitioning

- Delta data is stored as multiple small files when loaded.
- Delta Partitioning divides those data file into useful slices so they can be retrieved quickly and effectively.

Delta Partitioning

```
# Define DataFrame over root entity folder
( df.write
    .format("delta")
    .mode("overwrite")
    .partitionBy("<ColumnName>")
    .save("deltaFilePath")
)
```

To write delta partition you use the **.partitionBy()** command

No Partitioning

	path	name	size
1	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Exported/_delta_log/	_delta_log/	0
2	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Exported/part-00000-829f8ebc-6be6-4491-a1ba-f9bb93311d5f-c000 snappy.parquet	part-00000-829f8ebc-6be6-4491-a1ba-f9bb93311d5f-c000 snappy.parquet	161807024
3	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Exported/part-00001-a95fa1fc-a3c5-4249-b09b-a2e475442757-c000 snappy.parquet	part-00001-a95fa1fc-a3c5-4249-b09b-a2e475442757-c000 snappy.parquet	154872635
4	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Exported/part-00002-459fd8dc-f4ce-4dcf-bfd9-d87bd2f9c25-c000 snappy.parquet	part-00002-459fd8dc-f4ce-4dcf-bfd9-d87bd2f9c25-c000 snappy.parquet	153406419
5	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Exported/part-00003-2f681144-3098-4adf-a917-b5800c8557d5-c000 snappy.parquet	part-00003-2f681144-3098-4adf-a917-b5800c8557d5-c000 snappy.parquet	154855624
6	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Exported/part-00004-e30f11fb-008d-4029-87bd-7953f9346e20-c000 snappy.parquet	part-00004-e30f11fb-008d-4029-87bd-7953f9346e20-c000 snappy.parquet	147625269
-	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Exported/part-00005-b470a4e3-a8fb-413e-85c4-320e8d82ca3e-	part-00005-b470a4e3-a8fb-413e-85c4-320e8d82ca3e-c000 snappy.parquet	144943806

Partitioned

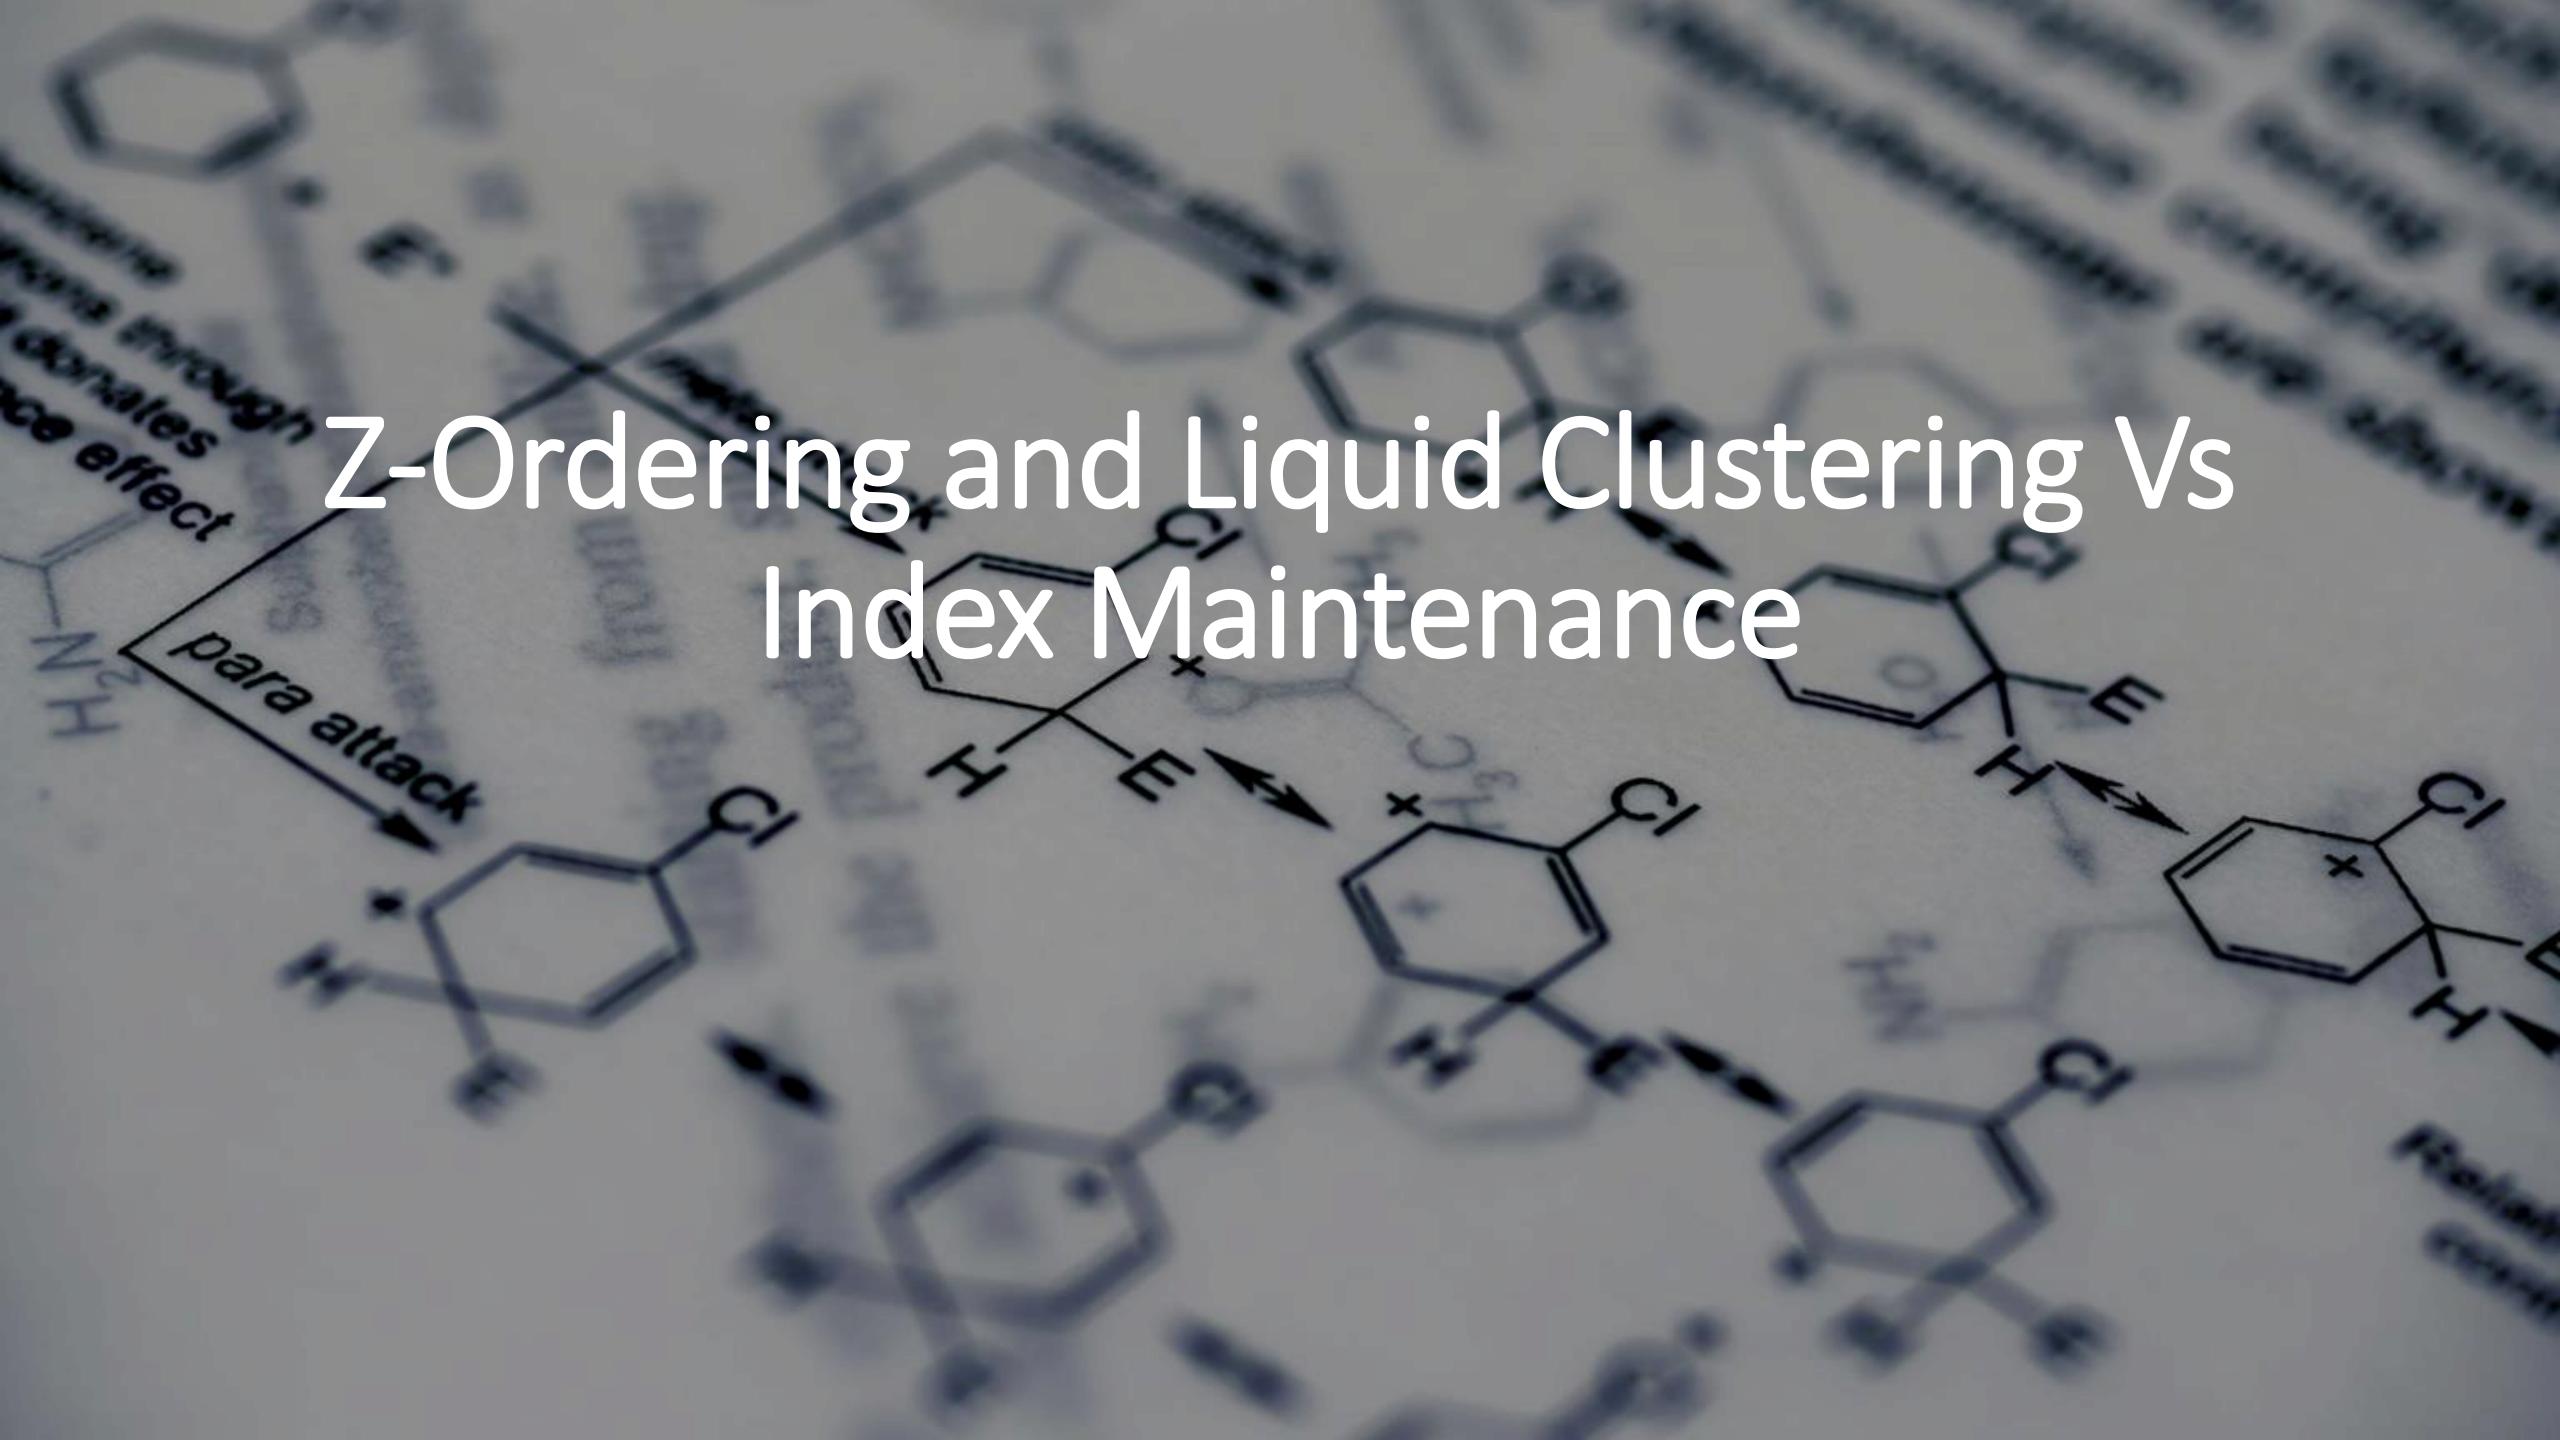
	path	name	size
1	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Partitioned/PULocationID=1/	PULocationID=1/	0
2	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Partitioned/PULocationID=10/	PULocationID=10/	0
3	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Partitioned/PULocationID=100/	PULocationID=100/	0
4	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Partitioned/PULocationID=101/	PULocationID=101/	0
5	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Partitioned/PULocationID=102/	PULocationID=102/	0
6	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Partitioned/PULocationID=104/	PULocationID=104/	0
7	dbfs:/mnt/deltalake/MasteringDelta/NYCTaxi/Yellow_TripData/Partitioned/PULocationID=105/	PULocationID=105/	0

Delta Demo

- Partitioning



Z-Ordering and Liquid Clustering Vs Index Maintenance



What problem are we solving?

- Improve performance when reading data
- We don't want to have to scan the entire dataset to find rows that match our query
- Ordering data on a granular row by row level for more optimal reading



z-Ordering in Delta

Z-Ordering

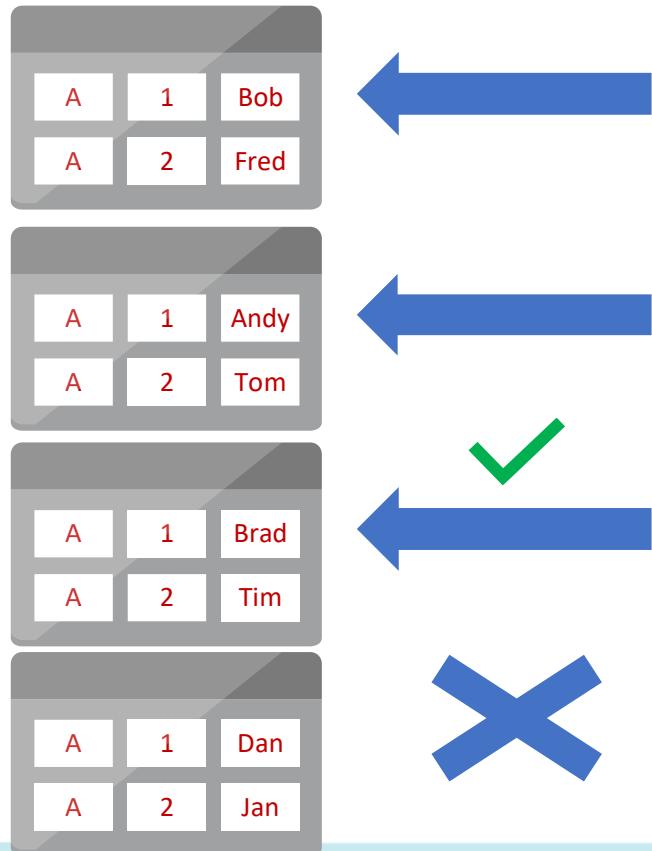
“Sort the data on specific columns before writing to files, to optimize data skipping”

```
--Optimize an entire table  
OPTIMIZE [database].[table] ZORDER BY [ColumnName]
```

Z-Order can be expensive, as with any sort-based Operation. It is sometimes better to perform this as an out-of-hours maintenance operation



Z-Ordering Explained



```
SELECT count(*) FROM Employees  
WHERE Name = 'Brad'
```

- The small files are not ordered
- SQL statement to query data
- Data skipping will look at all files until it finds what it's looking for

Z-Ordering Explained

A	1	Bob
A	2	Fred
A	1	Andy
A	2	Tom
A	1	Brad
A	2	Tim
A	1	Dan
A	2	Jan

OPTIMZE

A	1	Andy
A	1	Bob
A	1	Brad
A	2	Dan

**SELECT count(*)
FROM Employees
WHERE Name = 'Brad'**

ZORDER BY
Name

A	2	Fred
A	1	Jan
A	2	Tim
A	2	Tom

Z-Order is similar to a clustered index - organising your data on disk to maximise queries for specific columns

Demo

- ZOrdering

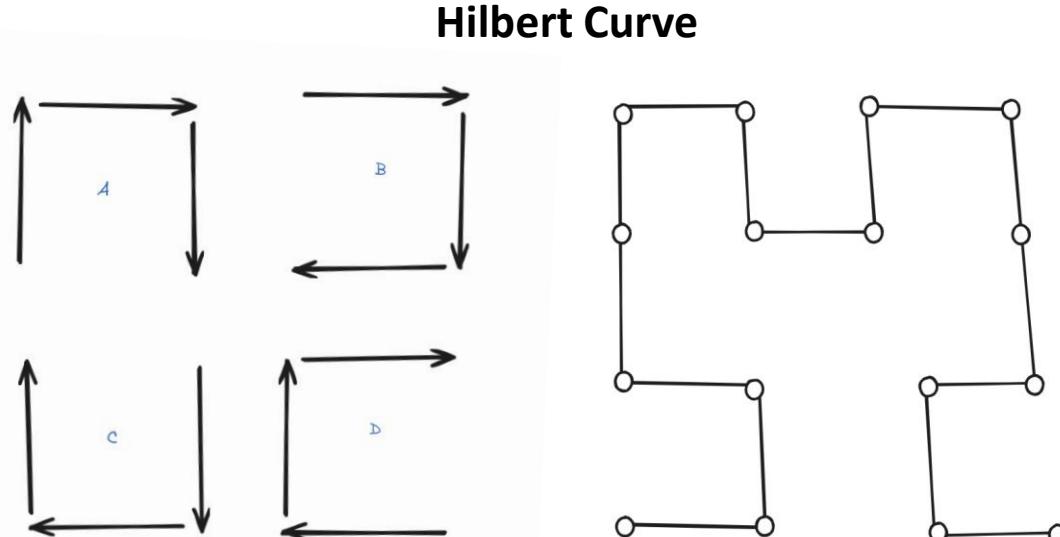
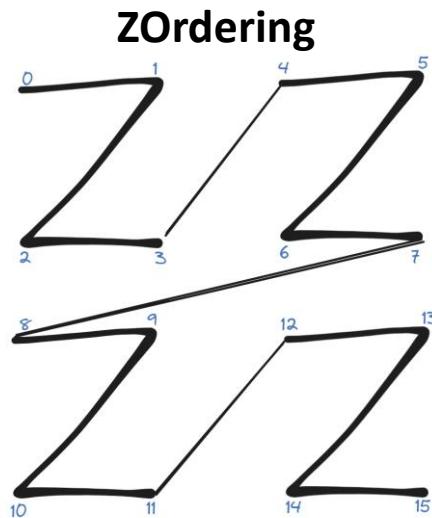


Databricks Liquid Clustering

Provides flexibility to redefine clustering columns without rewriting existing data, allowing data layout to evolve alongside analytic needs over time.

Hilbert Curve

The Hilbert curve is a way of mapping the multidimensional space into the one-dimensional space



How to Implement

```
CREATE TABLE table1(col0 int, col1 string) USING DELTA CLUSTER BY (col0);  
  
CREATE EXTERNAL TABLE table2 CLUSTER BY (col0) LOCATION 'table_location'  
AS SELECT * FROM table1;
```

Indexing in SQL Server

Z-Ordering is Index Maintenance

The order of the data is set by the Clustered Index

Index Maintenance

- Rebuilds
- Reorganization



Z-Ordering is Index Maintenance

Lets look at bit deeper into what is happening and why this is needed.

We have two different problem for row-store and column-store indexes as they used the two different structures

Lets start with Row based!

Z-Ordering is Index Maintenance

When new records are added to the tables, the indexes are updated!

Lets keep with Person.Person

We have a B-Tree with and see what can happen when we make some changes!

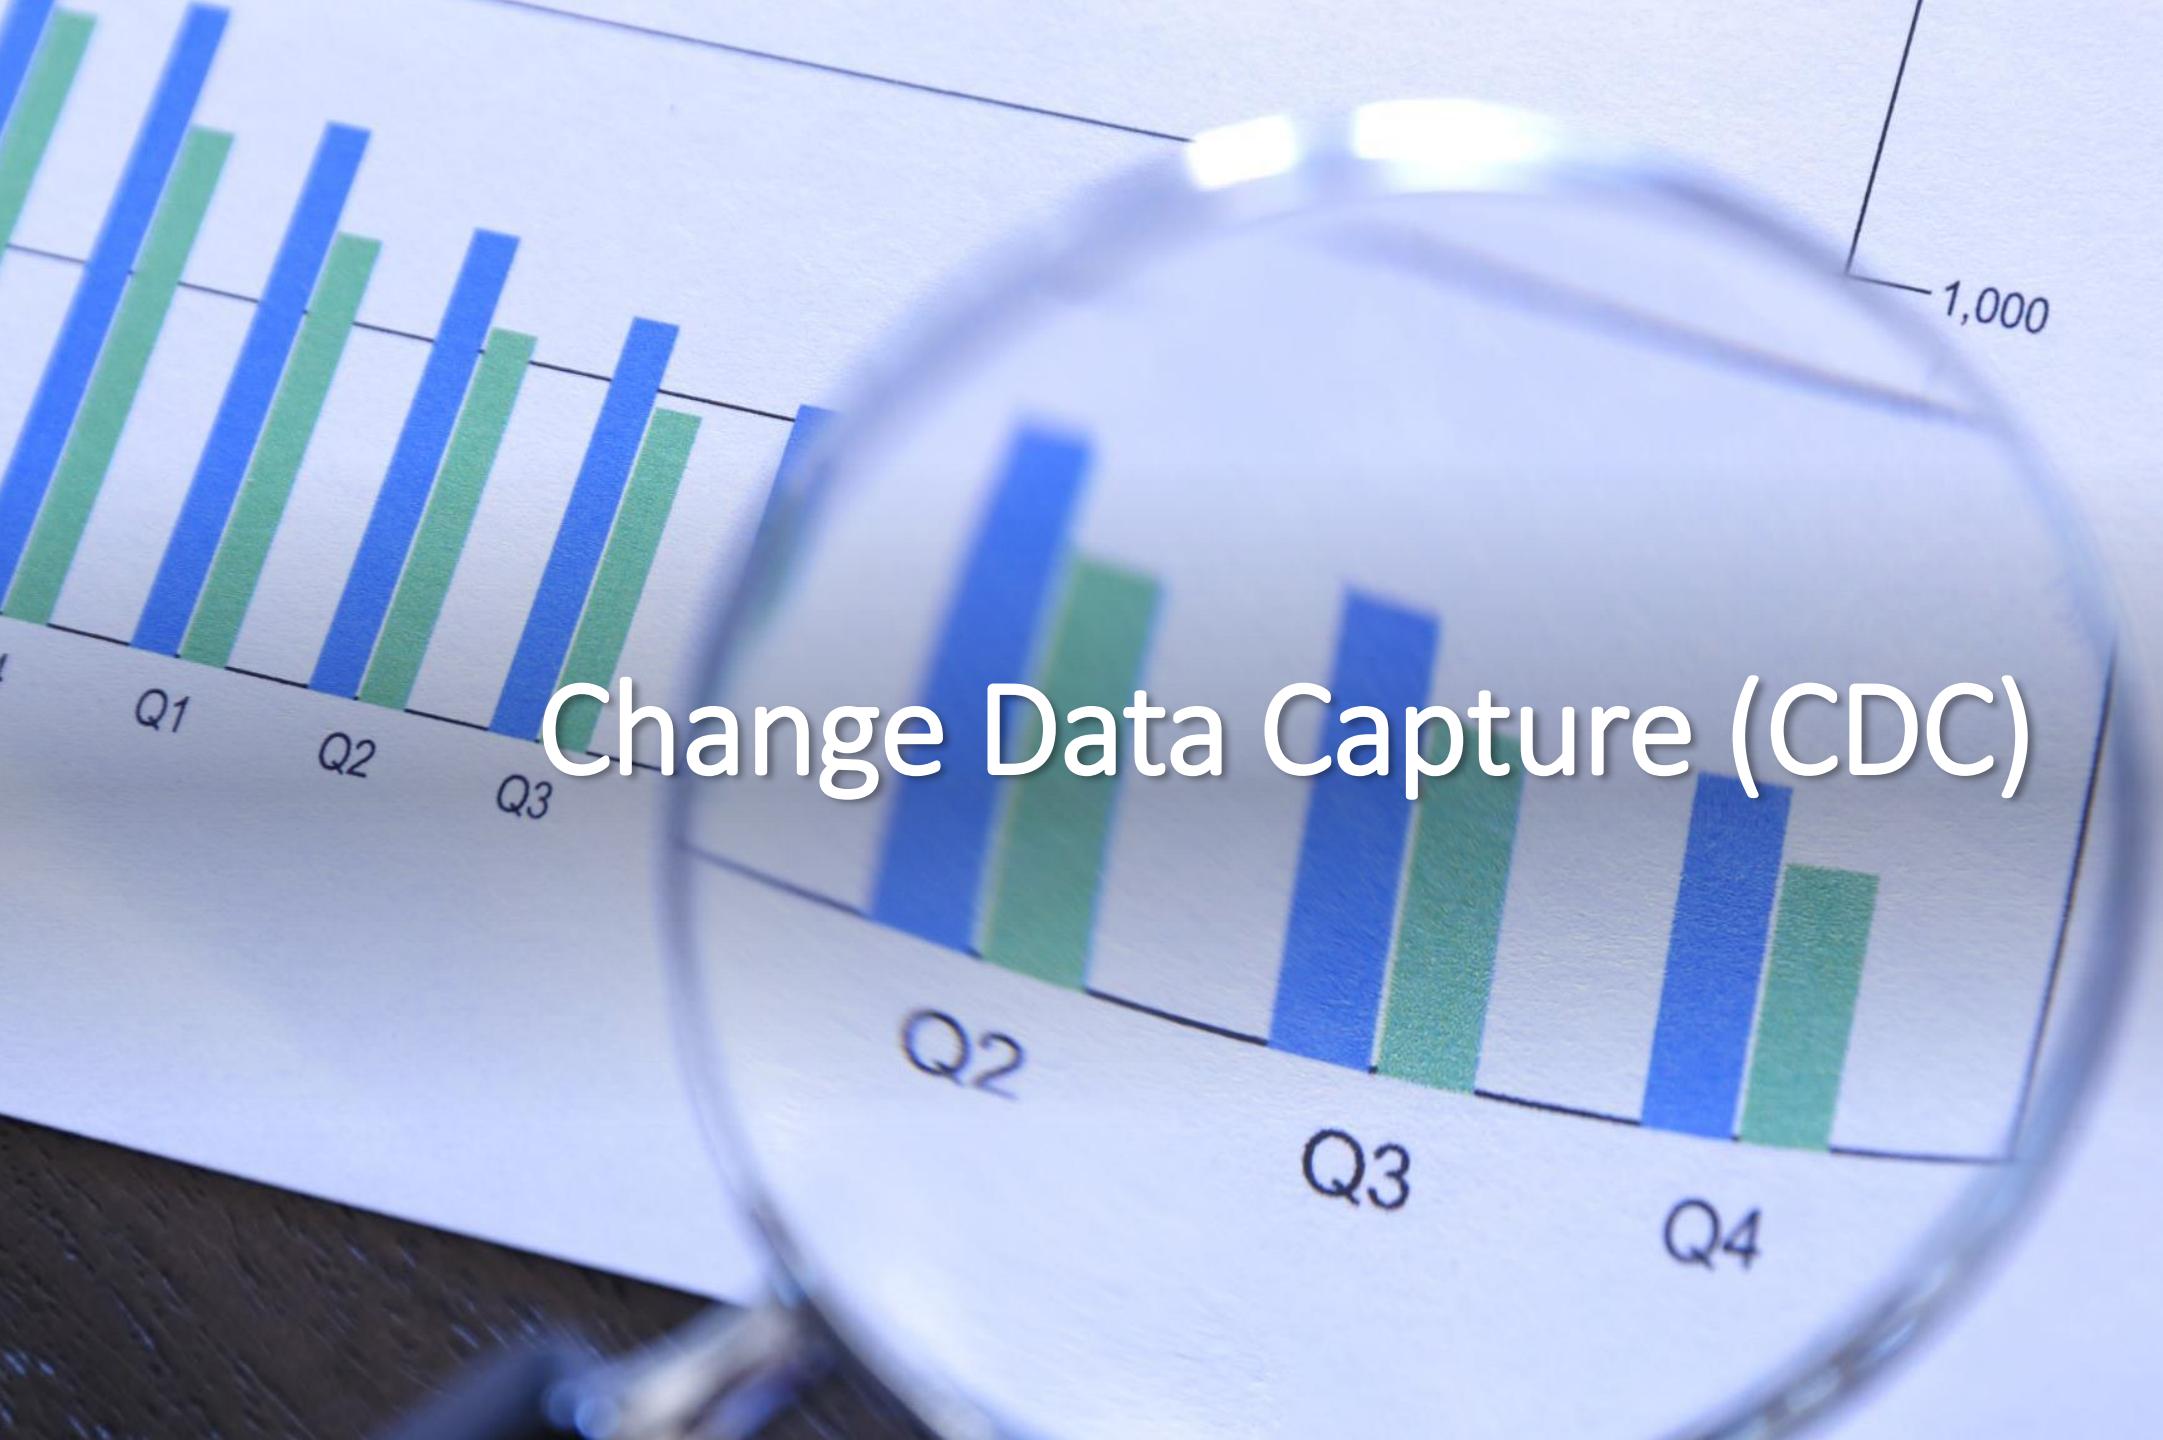
Z-Ordering is Index Maintenance

Column-store works in works in a different way.



1,000

Change Data Capture (CDC)



The only constant is Change!

What problem are we solving?

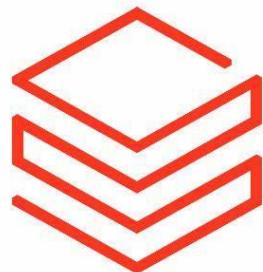
Can you tell me what has changed?



Change Data Capture
Change Tracking*

Can you tell me what this look like at this time?

Temporal Tables



Change Data
Capture

Delta Time Travel



Change Data Capture in SQL Server

Change Data Capture

We can track changes of tables by enabling CDC at a database level.

We then need to add each table.

Once this is done, we can see what has change and what it was previously.

Let take a look at Person.Person

Retention, These can get big, Log and table space!

Change Data Capture

```
/*Enable on the database */
```

```
EXEC sys.sp_cdc_enable_db  
GO
```

```
/*Enable on the table, Requires a PK */
```

```
EXEC sys.sp_cdc_enable_table  
    @source_schema = N'Person',  
    @source_name    = N'PersonCDC',  
    @role_name      = NULL,  
    @supports_net_changes = 1  
GO
```

Change Data Capture

```
/* Select with cdc function */

SELECT[__$start_lsn]
    ,[__$end_lsn]
    ,[__$seqval]
    ,[__$operation]
    ,[__$update_mask]
    ,[FirstName]
    ,[MiddleName]
    ,[LastName]
    ,[ModifiedDate]
    ,sys.fn_cdc_map_lsn_to_time($__start_lsn) AS 'CDCModifiedDate'
FROM [AdventurePlay].[cdc].[Person_PersonCDC_CT]
```

Change Data Capture

```
UPDATE [Person].[PersonCDC]
      SET [FirstName] = 'James',
          [MiddleName] = 'C',
          [LastName] = 'Yarrow'
     WHERE [BusinessEntityID] = 4000
```

	Results	Messages								
	\$start_lsn	\$end_lsn	\$seqval	\$operation	\$update_mask	FirstName	MiddleName	LastName	ModifiedDate	CDCModifiedDate
1	0x0000003100014E9A0003	NULL	0x0000003100014E9A0002	3	0x0070	Jeremiah	E	Martin	2014-02-02 00:00:00.000	2024-02-26 18:32:27.410
2	0x0000003100014E9A0003	NULL	0x0000003100014E9A0002	4	0x0070	James	C	Yarrow	2014-02-02 00:00:00.000	2024-02-26 18:32:27.410

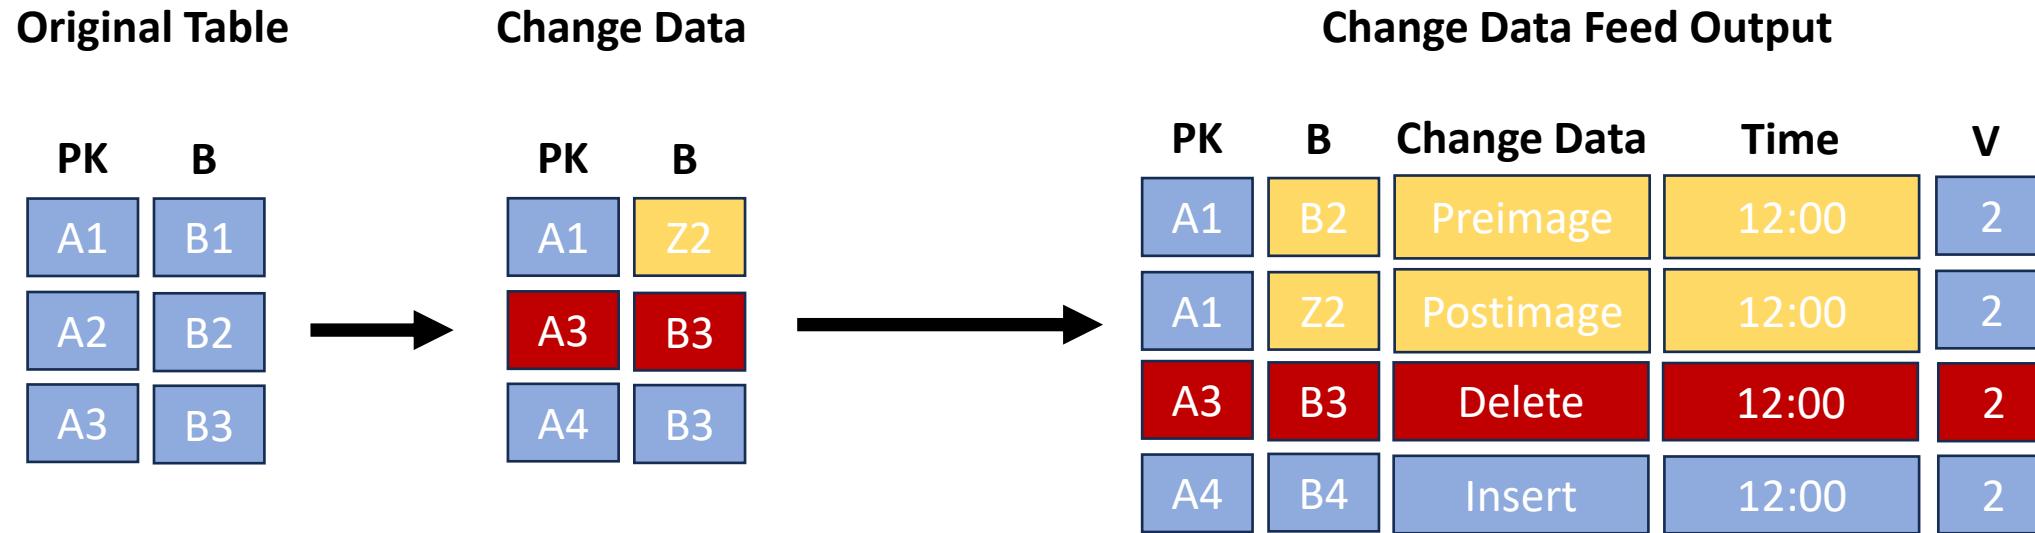
Change Data Capture in Databricks

Change Data Capture (CDC)

Databricks has a feature to handle CDC called Change Data Feed (CDF), which allows Databricks to track row-level changes between versions of a Delta table



How does it work



We have four change types:

- Preimage
- Postimage
- Delete
- Insert

How to Implement

```
CREATE TABLE student (id INT, name STRING, age INT) TBLPROPERTIES  
(delta.enableChangeDataFeed = true)
```

```
ALTER TABLE myDeltaTable SET TBLPROPERTIES (delta.enableChangeDataFeed =  
true)
```

```
set spark.databricks.delta.properties.defaults.enableChangeDataFeed = true;
```

How to Read

```
# version as ints or longs  
spark.read.format("delta")  
.option("readChangeFeed", "true")  
.option("startingVersion", 0)  
.option("endingVersion", 10)  
.table("myDeltaTable")
```

```
# timestamps as formatted timestamp  
spark.read.format("delta")  
.option("readChangeFeed", "true")  
.option("startingTimestamp", '2021-04-  
21 05:45:46')  
.option("endingTimestamp", '2021-05-  
21 12:00:00')  
.table("myDeltaTable")
```

Delta History and Time Travel vs Temporal Tables

What problem are we trying to solve?

Can you tell me what this look like at this time?

SQL Server Temporal Tables

Temporal Tables

- Allow you to keep history of data changes
- Point in time analysis
- Used for auditing
- Retention

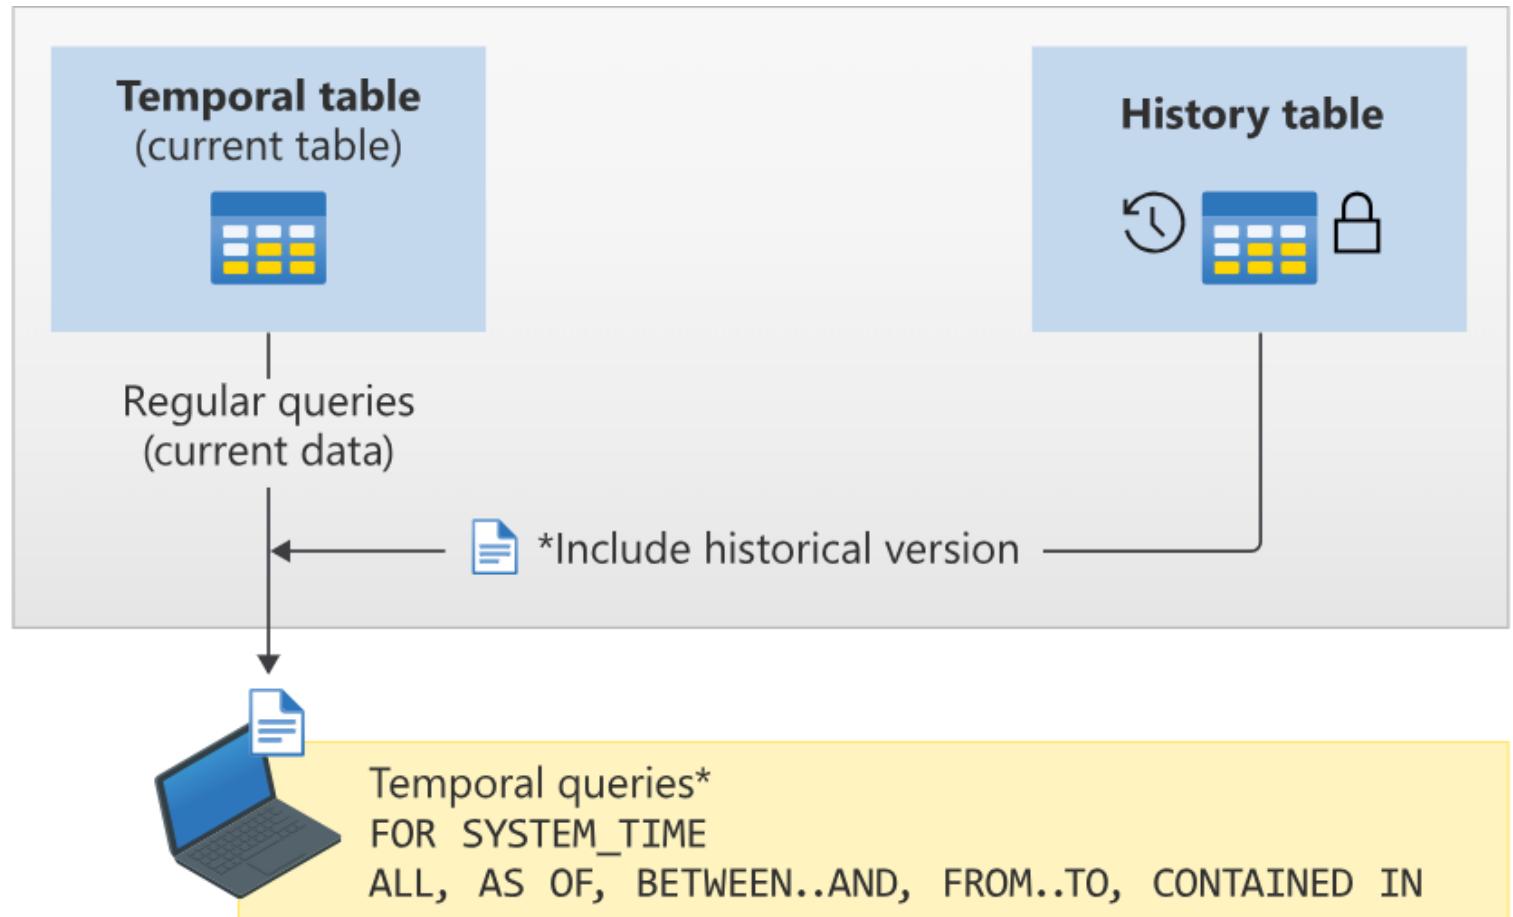


How do they work?

- Columns to define the period of validity
 - Period Start column
 - Period End Column
- History table



How can you query the data?



Demo





Welcome

DemoScripts.ipynb

SQLVsDeltaBricks.ipynb X

...

C: > Users > annet > OneDrive > Presentations > DeltaVsSQLServer > SQLVsDeltaBricks.ipynb

+ Cell ▾

Run all

Kernel SQL

Attach to NOODLE\BRANDS, Adventur

↑ ↓ ⚡ ⚡ ⚡ ⚡ ↗

To edit a current table

- Add in the 2 Period start and end columns
- State the history table



```
[ ] 1 Alter table Person.Person
2     Add
3         PeriodFrom datetime2 GENERATED ALWAYS AS ROW START NOT NULL
4             DEFAULT SYSUTCDATETIME()
5     , PeriodTo datetime2 GENERATED ALWAYS AS ROW END NOT NULL
6             DEFAULT CAST('9999-12-31 23:59:59.999999' AS datetime2),
7     PERIOD FOR SYSTEM_TIME (PeriodFrom,PeriodTo);
```



SQL

```
[ ] 1
2 Alter table Person.Person
3 set
4     (SYSTEM_VERSIONING = ON
5         (HISTORY_TABLE = Person.Person_history))
6
```



SQL

Now lets look at the data in the table to see what it looks like

PROBLEMS 2

OUTPUT

TERMINAL

TASKS

Tasks

▼

☰ 🔒 📁 ⌂ ⌃

⊗ 2 △ 0



Limitations of Temporal Tables

- History tables have to be in the same database as the source table
- You can only delete from temporal tables you can't truncate them
- You can't make any changes to the history table
- History tables can't have any constraints e.g. primary or foreign keys



Delta History and Time Travel in Databricks

Delta Time Travel

- Delta Lake allow you to query older version/snapshots of a Delta table using Time Travel.
- Delta Lake stored a 30-days version history by default.
- Useful for debugging, auditing, rolling back for data quality or to reproduce experiments.
- Can query an older versions using Python or SQL syntax.

Delta Time Travel

```
-- SQL syntax
SELECT count(*)
FROM <databaseName>.<tableName>
VERSION AS OF 0
```

```
-- DataFrameReader options
( spark.read
    .format("delta")
    .option("versionAsOf", 0)
```

	version	timestamp
4	4	2021-08-09T10:48:28.000+0000
5	3	2021-08-09T10:47:53.000+0000
6	2	2021-08-09T10:45:52.000+0000
7	1	2021-08-09T10:45:49.000+0000
8	0	2021-08-09T10:45:29.000+0000

Delta Time Travel Demo

- History
- Version

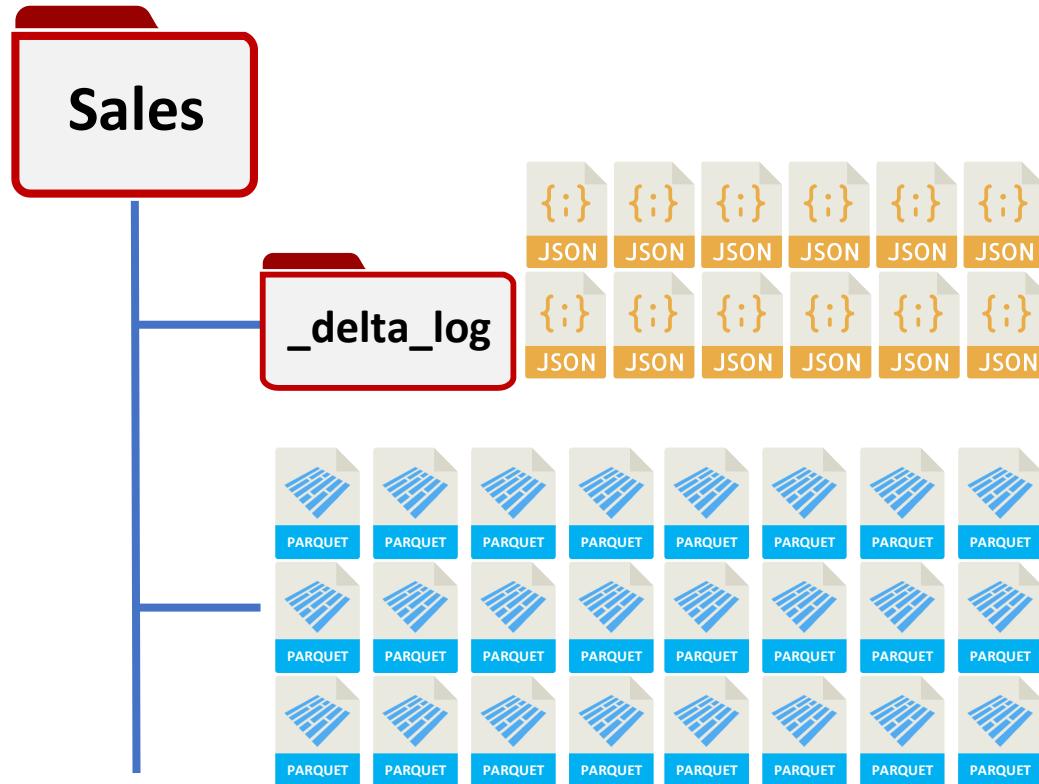
Delta Vacuum vs SQL Transactions

What problem are we trying to solve?

Clean up old data

Delta Vacuum

Obsolete Files



Removing historical To remove
obsolete history files, Delta
has the **VACUUM** command

This command physically deletes data
files older than a specified date

You **CANNOT** time travel past dates
where history has been vacuumed

Using the Vacuum Command

Vacuuming in SQL:

```
--Vacuum Table using defaults
VACUUM [database].[table]

--Vacuum using path not Hive table
VACUUM '/mnt/lake/BASE/myTable/'

--VACUUM for a non-default time period
VACUUM [database].[table] RETAIN 168 HOURS

--TEST THE VACUUM BEFORE YOU RUN IT
VACUUM [database].[table] RETAIN 168 HOURS DRY RUN
```

Using the python
DeltaTable object:

```
# Vacuum Table using defaults
deltaTable.vacuum()

# Vacuum Table for files older than 7 days (168 hours)
deltaTable.vacuum(168)
```



Demo

- Vacuum



SQL Server Retention

Data Retention

- Audit Logs
- Temporal Tables
- Backups
- Change Data Capture



Primary/Foreign Keys and Constraints

What problem are we solving?

- How can we keep referential integrity?
- How can we join data together across data?



DELTA Primary & Foreign Keys

What is Databricks Unity Catalog?

Unity Catalog provides centralized access control, auditing, lineage, and data discovery capabilities across Databricks workspaces.

- Define once, secure everywhere
- Standards-compliant security model
- Built-in auditing and lineage
- Data discovery
- System tables



Governance Layer



Delta Constraints - Unity Catalog

- Primary Keys
- Foreign Keys
- Identity Columns
- Not Null
- Unique
- Check

Delta Constraints - Databricks Only

Primary Keys / Foreign Keys / Identity Columns (Surrogate Keys)

```
# Primary Key
CREATE TABLE <tableName> (
    <col1> INTEGER NOT NULL,
    <col2> INTEGER NOT NULL,
    CONSTRAINT <name> PRIMARY KEY(<col1>, <col2>)
);
```

```
# Foreign Key
CREATE TABLE <tableName> (
    <pkCol> INTEGER NOT NULL PRIMARY KEY,
    <col1> INTEGER,
    <col2> INTEGER,
    CONSTRAINT <name> FOREIGN KEY(<col1>, <col2>)
    REFERENCES <refName>
);
```

```
# Identity Columns
CREATE OR REPLACE TABLE <tableName> (
    <pkCol> BIGINT GENERATED ALWAYS AS IDENTITY (START WITH 1 INCREMENT BY 1) PRIMARY KEY,
    <col1>, <col2>);
```


SQL Primary and Foreign keys

Primary and Foreign Keys

SQL Server Primary Key

- One per table and Always Unique
- At least one Column
- Not Null
- Just a constraint
- Can be extracted for your data checks down the pipeline

SQL Server will often create your cluster index on your PK



Databricks Security vs SQL Security

SQL Server Security

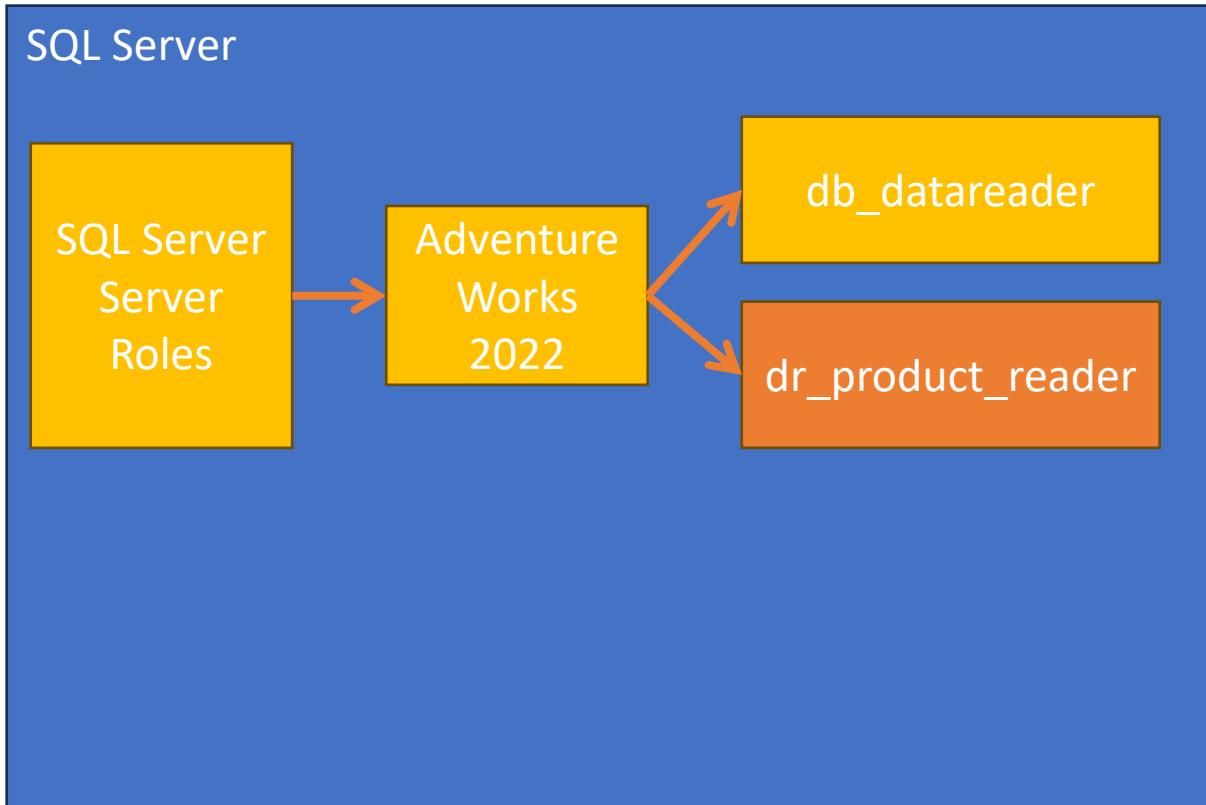
SQL Server Security

- All (Most) Data is stored within the database
- Access is controlled via SQL Server
- SQL Server can be its only authority with SQL Logins
- SQL Server holds the file locks at all times when started



SQL Server Security

- SQL Server have a server within this databases
- Server Roles are applied, then database roles are applied if needed
- Database roles have built-in Roles



SQL Server Security

We can map ~~Azure Active Directory~~ Entra ID, First we map to a Login at the Server level.

```
USE [master]
GO
--CREATE LOGIN [EntraIDName] FROM EXTERNAL PROVIDER
CREATE LOGIN [MaskingTestUser] WITHOUT LOGIN
```

We then need to create a User at the Database level.

```
USE [AdventureWorks2022]
GO
CREATE USER [MaskingTestUser] FOR LOGIN [MaskingTestUser]
GO
ALTER ROLE [dr_sales_reader] ADD MEMBER [MaskingTestUser]
GO
```

SQL Server Security

There are built in roles that can grant read across the database [db_reader] but lets limit this to the sales schema

```
USE [AdventureWorks2022]
GO
CREATE ROLE [dr_sales_reader]
GO
GRANT SELECT ON SCHEMA::[Sales] TO [dr_sales_reader]
GO
```

SQL Server Security

We can then test this with the user we created.

```
EXECUTE AS USER = 'MaskingTestUser';
SELECT TOP (10) *
FROM [Sales].[CreditCard];
SELECT TOP (10) *
FROM [Person].[Person];
REVERT;
```

	CreditCardID	CardType	CardNumber	ExpMonth	ExpYear	ModifiedDate
1	1	SuperiorCard	33332664695310	11	2006	2013-07-29 00:00:00.000
2	2	Distinguish	55552127249722	8	2005	2013-12-05 00:00:00.000
3	3	ColonialVoice	77778344838353	7	2005	2014-01-14 00:00:00.000
4	4	ColonialVoice	77774915718248	7	2006	2013-05-20 00:00:00.000
5	5	Vista	11114404600042	4	2005	2013-02-01 00:00:00.000
6	6	Distinguish	55557132036181	9	2006	2014-04-10 00:00:00.000
7	7	Distinguish	55553635401028	6	2007	2013-02-01 00:00:00.000
8	8	SuperiorCard	33336081193101	7	2007	2013-06-30 00:00:00.000
9	9	Distinguish	55553465625901	2	2005	2013-09-23 00:00:00.000
10	10	SuperiorCard	33332126386493	8	2008	2011-08-31 00:00:00.000

```
Msg 229, Level 14, State 5, Line 15
The SELECT permission was denied on the object 'Person',
database 'AdventureWorks2022', schema 'Person'.
```

SQL Server Security

We can also apply data masking at the column level.

```
ALTER TABLE [Sales].[CreditCard]
ALTER COLUMN [CardNumber] ADD MASKED WITH (FUNCTION = 'default()');
```

This can be tested with a local user for test which can also be put in your pipeline!

```
EXECUTE AS USER = 'MaskingTestUser';
SELECT TOP (10) *
FROM [Sales].[CreditCard];
REVERT;
```

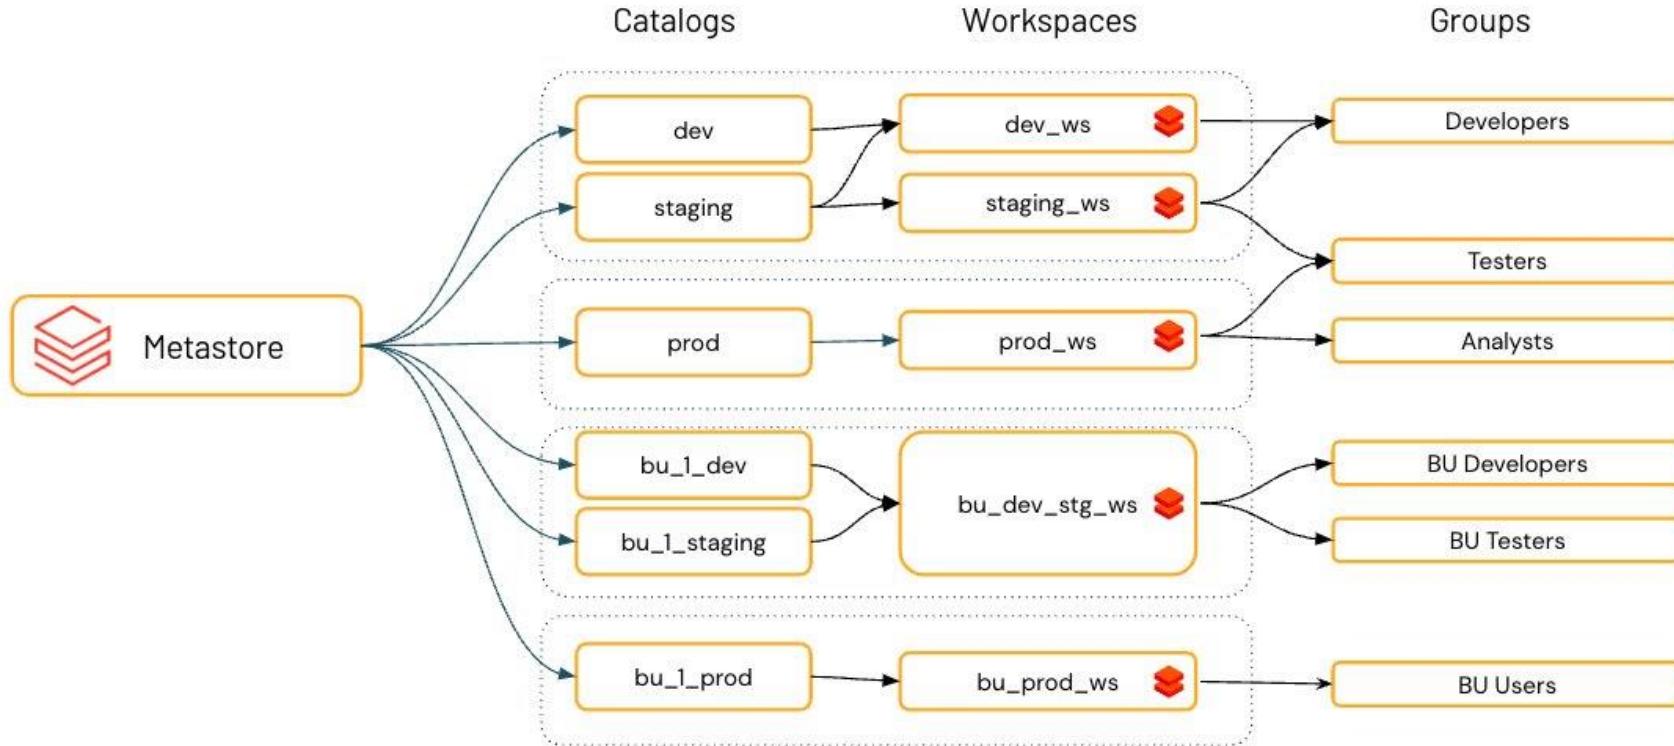
	CreditCardID	CardType	CardNumber	ExpMonth	ExpYear	ModifiedDate
1	1	SuperiorCard	xxxx	11	2006	2013-07-29 00:00:00.000
2	2	Distinguish	xxxx	8	2005	2013-12-05 00:00:00.000
3	3	ColonialVoice	xxxx	7	2005	2014-01-14 00:00:00.000
4	4	ColonialVoice	xxxx	7	2006	2013-05-20 00:00:00.000
5	5	Vista	xxxx	4	2005	2013-02-01 00:00:00.000
6	6	Distinguish	xxxx	9	2006	2014-04-10 00:00:00.000
7	7	Distinguish	xxxx	6	2007	2013-02-01 00:00:00.000
8	8	SuperiorCard	xxxx	7	2007	2013-06-30 00:00:00.000
9	9	Distinguish	xxxx	2	2005	2013-09-23 00:00:00.000
10	10	SuperiorCard	xxxx	8	2008	2011-08-31 00:00:00.000

Databricks (UC) Security

Previously

- Secured Data on ADLS instance via Access Control Lists (ACLs)
- Couldn't sync AD (Active Directory) groups to Databricks

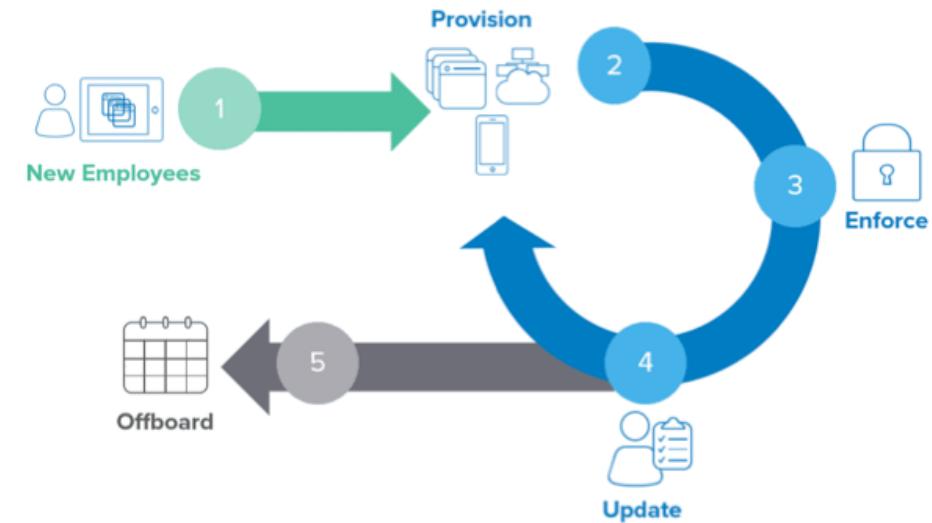
Security in Unity Catalog



SCIM Connector

System for Cross domain Security management

We can do this Using Microsoft Entra ID
(formerly Azure Active Directory) when
working with Unity Catalog



UC Security Examples

```
-- Let's grant all users a SELECT
```

```
GRANT SELECT ON TABLE customers TO `account users`;
```

```
-- We'll grant an extra MODIFY to our Data Engineer
```

```
-- Note: make sure you created the dataengineers group first as an account admin!
```

```
GRANT SELECT, MODIFY ON TABLE customers TO `dataengineers`;
```

UC Security Examples

```
-- get the current user (for informational purposes)
```

```
SELECT current_user(), is_account_group_member('account users');
```

```
-- Apply access rule to customers table.
```

```
-- country will be the column sent as parameter to our SQL function (region_param)
```

```
ALTER TABLE customers SET ROW FILTER region_filter ON (country);
```



UC Security Examples

```
CREATE OR REPLACE FUNCTION region_filter(region_param STRING)
RETURN
    is_account_group_member('bu_admin') or -- bu_admin can access all regions
    region_param like "US%" ;           -- non bu_admin's can only access regions
                                         containing US

-- Grant access to all users to the function for the demo by making all account users
owners. Note: Don't do this in production!
GRANT ALL PRIVILEGES ON FUNCTION region_filter TO `account users`;
```

Summary

- We have seen the modern challenges to data platforms
- We have different approaches to Compute and Storage and the additional complexity of Databricks
- We can have ACID transaction and Partitioning in both
- Both have Ordering data on a granular row by row level for more optimal reading
- Both have point in time querying and changed data capture
- We have a lot of overlap these days, but this also means they can play well
- Similarities in security approaches but still very different components to secure
- We have different advantages but also limitations

Thank you all!



<https://sqlb.it/?10976>

<https://github.com/AnnaWykes/delta-dbx-for-dba>

Please fill out the online feedback!

This has dropped off in our completions rate, since we have switched away for paper feedback.

Don't make us kill the trees, Please do the online feedback!